Terminologies

- Cluster computing -> is the kind of architecture that group of computers use their resources to perform the particular task by means of interconnected with them via network
- Distributed Computing
- Distributed Storage
- Auto Scaling
- Difference between Horizontal and Vertical Scaling

Cluster

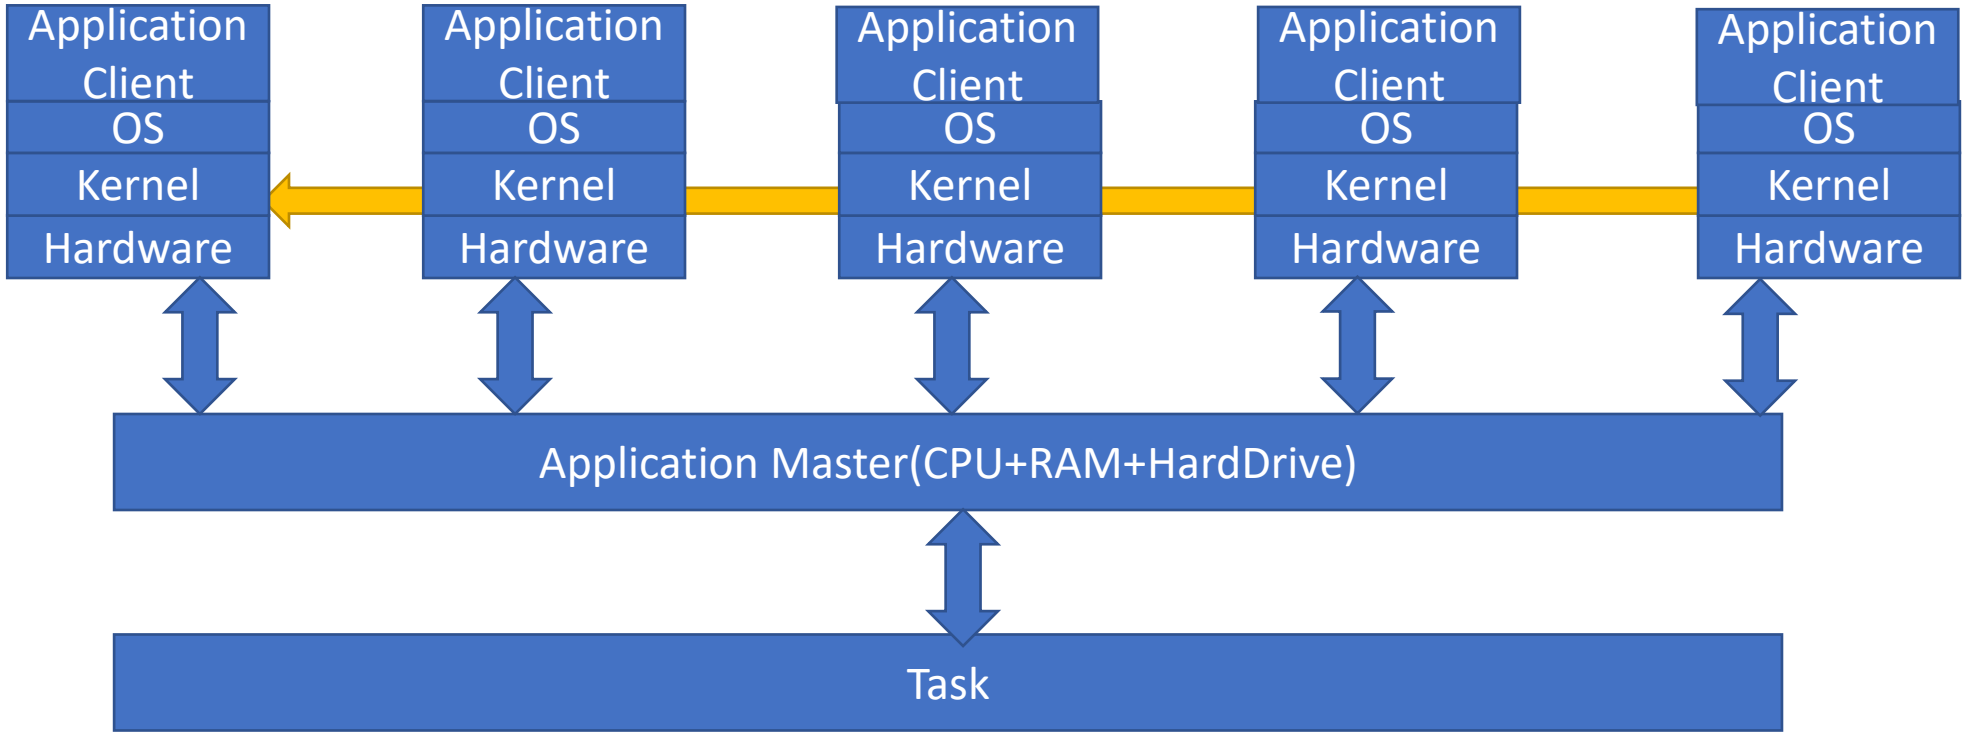
Application Client
OS
Kernel
Hardware
Application Client
OS
Kernel
Hardware
Application Client
OS
Kernel
Hardware
Application Client
OS
Kernel
Hardware
Application Client
OS
Kernel
Hardware
Application Master(CPU+RAM+HardDrive)
Task

Distributed Storage

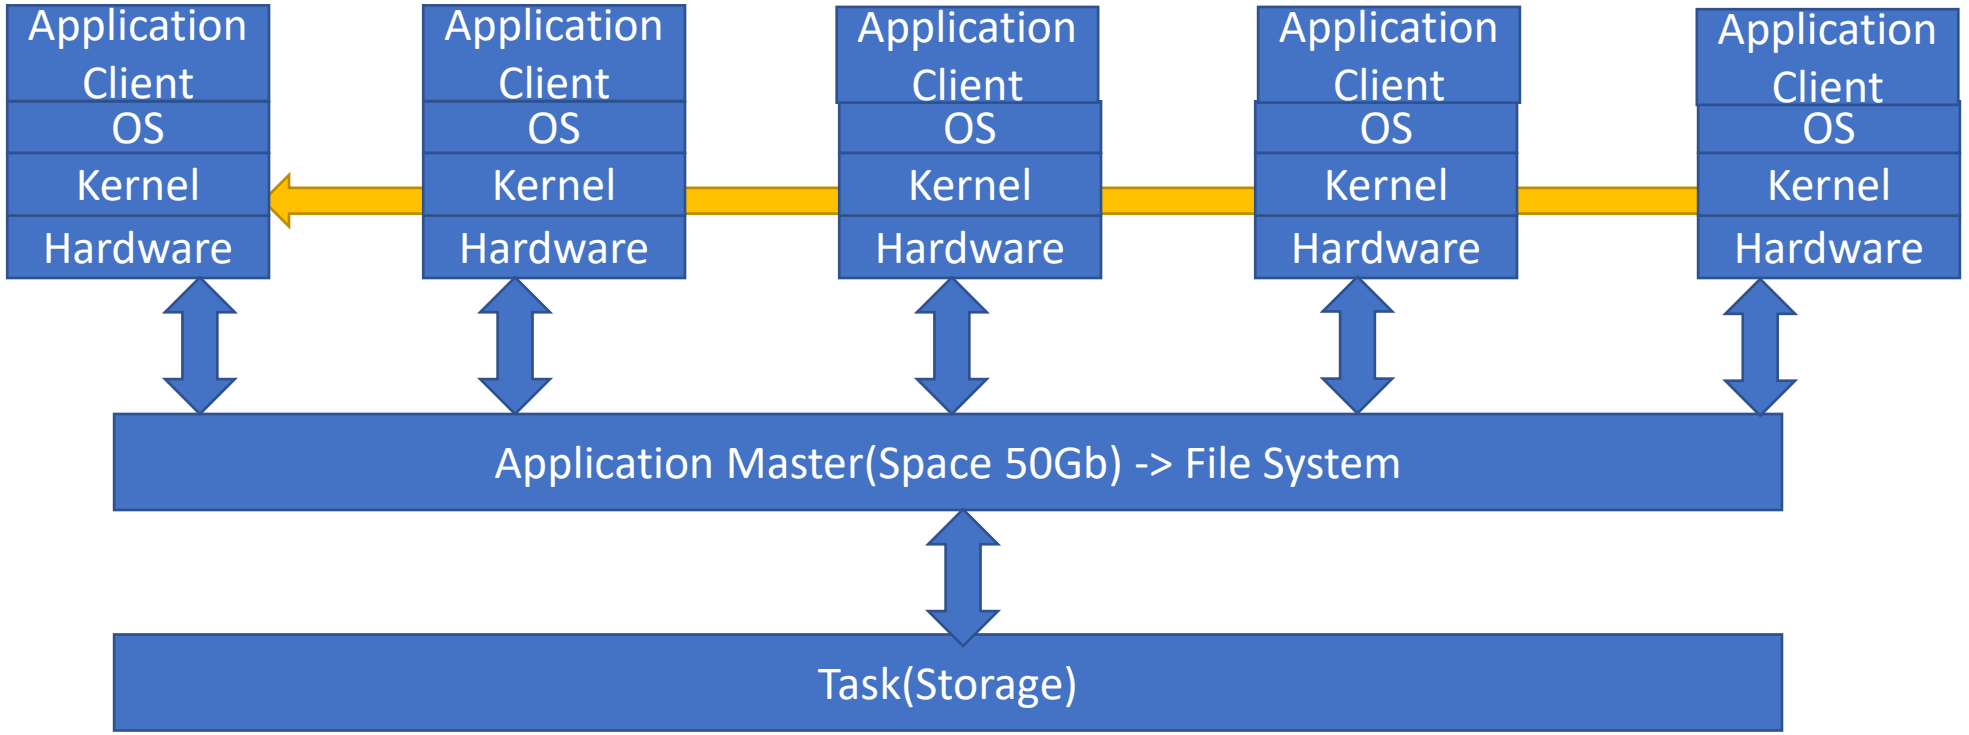

Distributed Computing

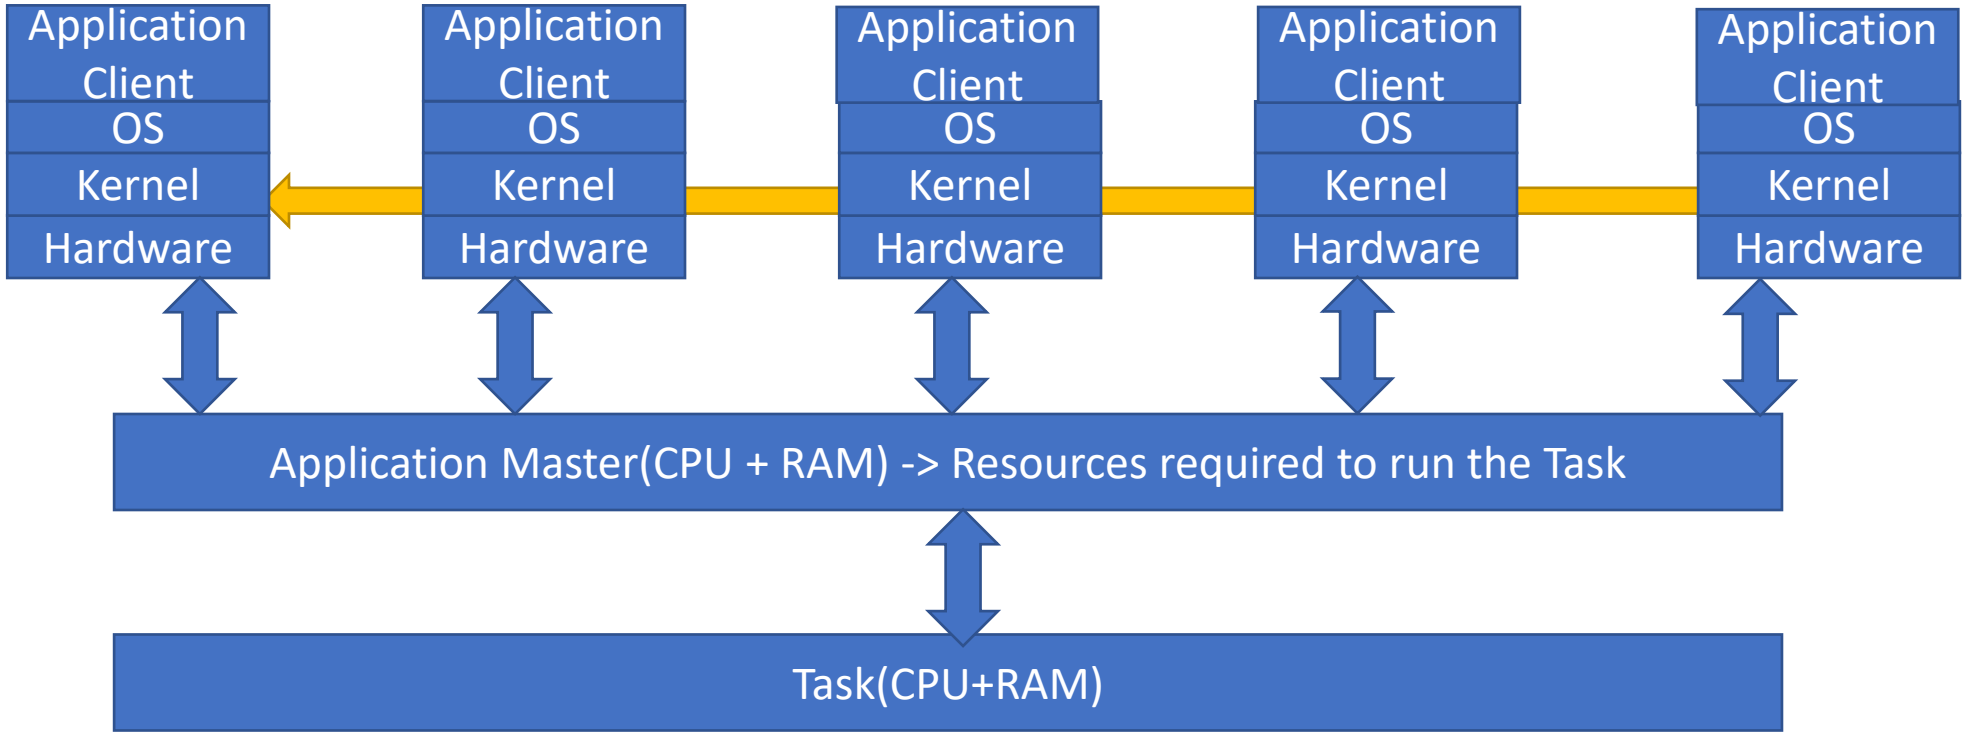

Auto Scaling

MYSQL Data Base

12th GEN i9(8 Core)
8GB Ram
1TB SSD

Vertical Scaling

While Scaling UP/Down ->
Shutdown -> Remove the Hardware

Application Downtime

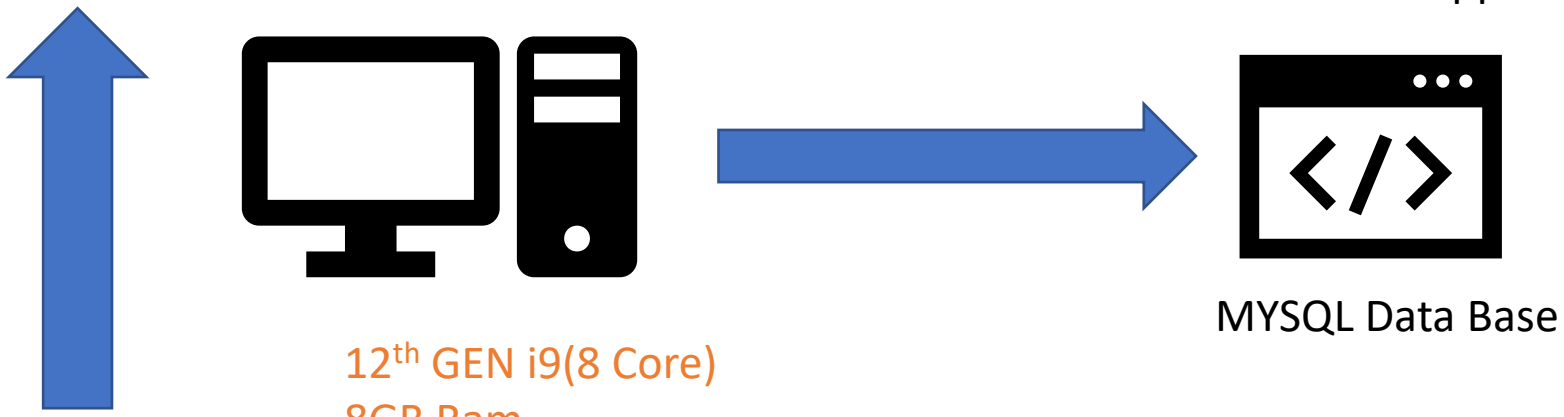

12th GEN i9(8 Core)
8GB Ram
1TB SSD

MYSQL Data Base

Horizontal Scaling

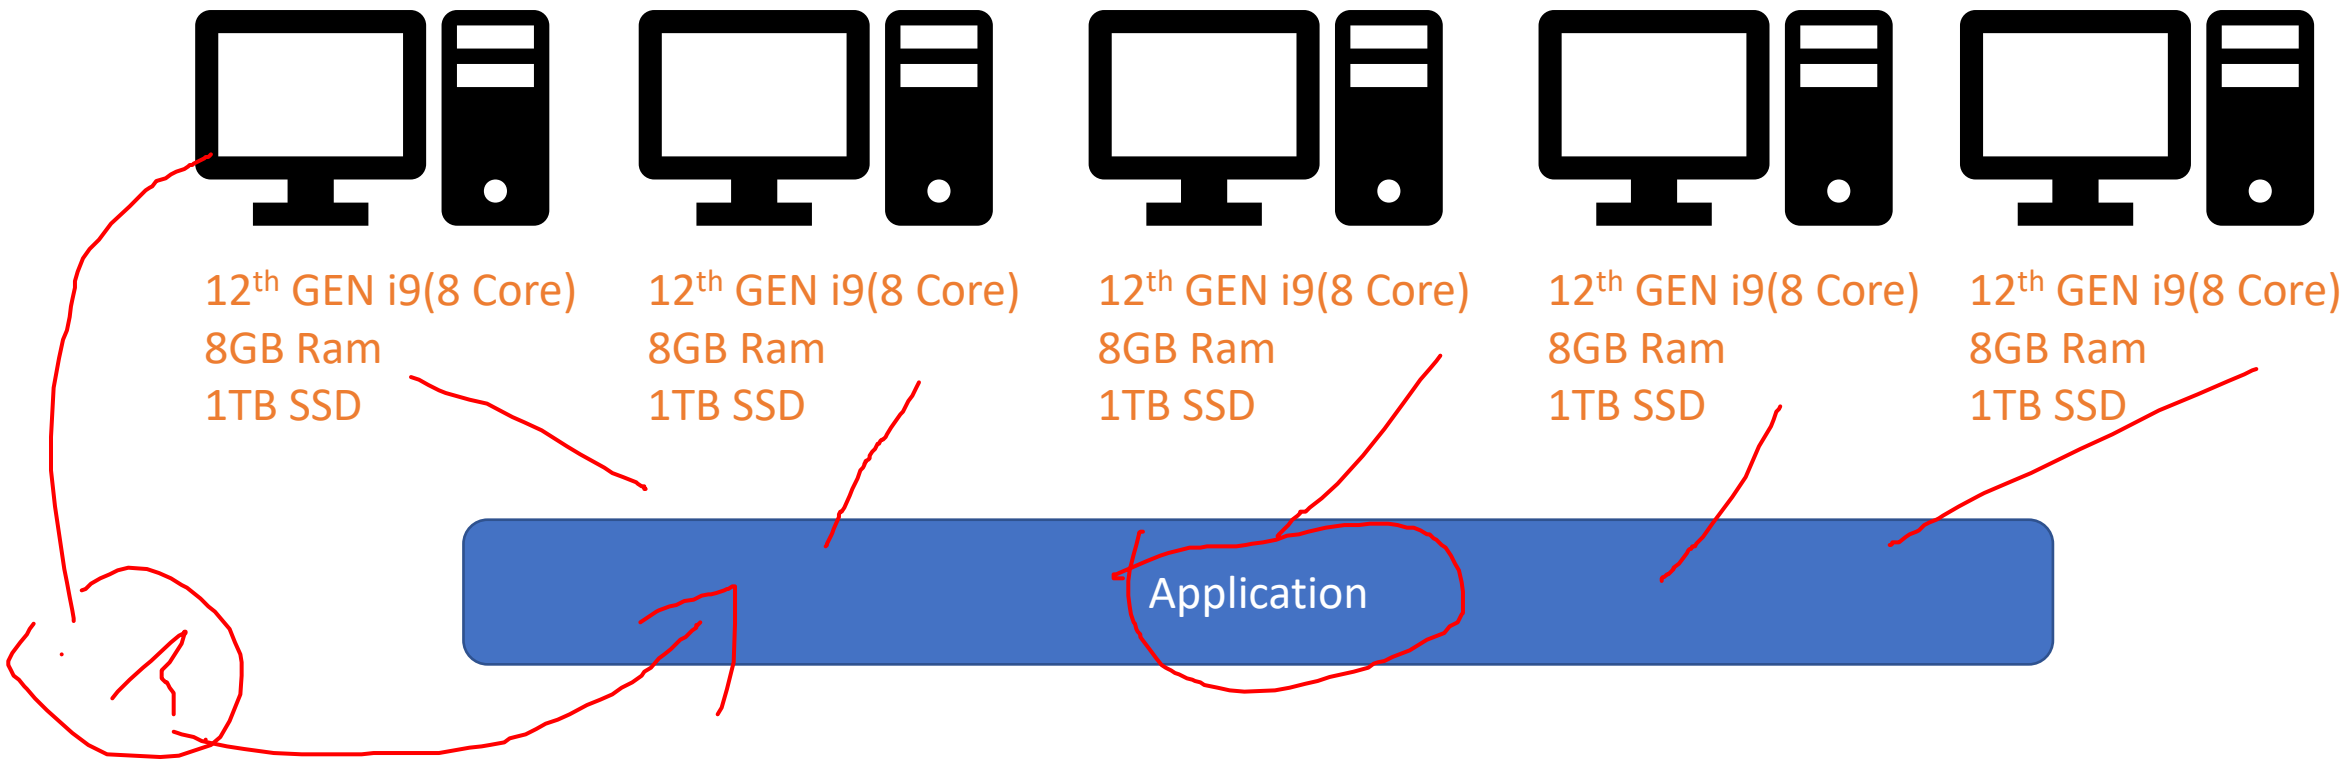

Bigdata

- Big data refers to data sets that are too large or complex to be dealt with by traditional data-processing application software(database).
- Volume -> doesn’t have any kind of limitation in handling the data volume(Distributed Storage)
- Velocity -> time taken to process the amount of data(Computational limitation) -> Distributed Computing
- Variety -> We don’t need to deal the files with ETL Concept -> As the Flat File -> it got variety of tools to deal with the data as same format without doing any kind of type conversion.

- 3 v
- Volume -> Storage
- Velocity -> Computation
- Variety ->

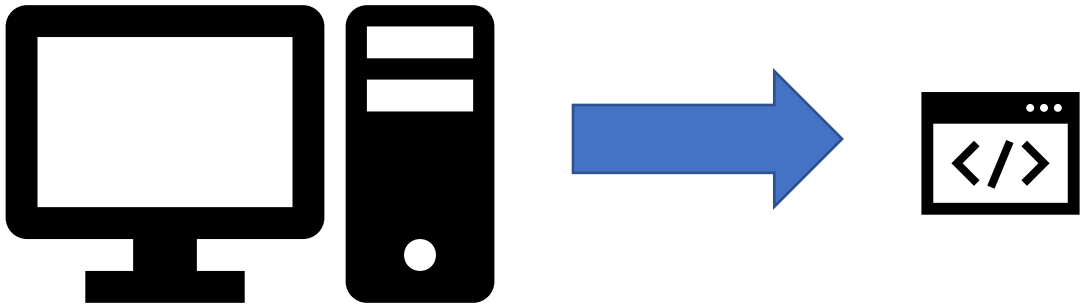

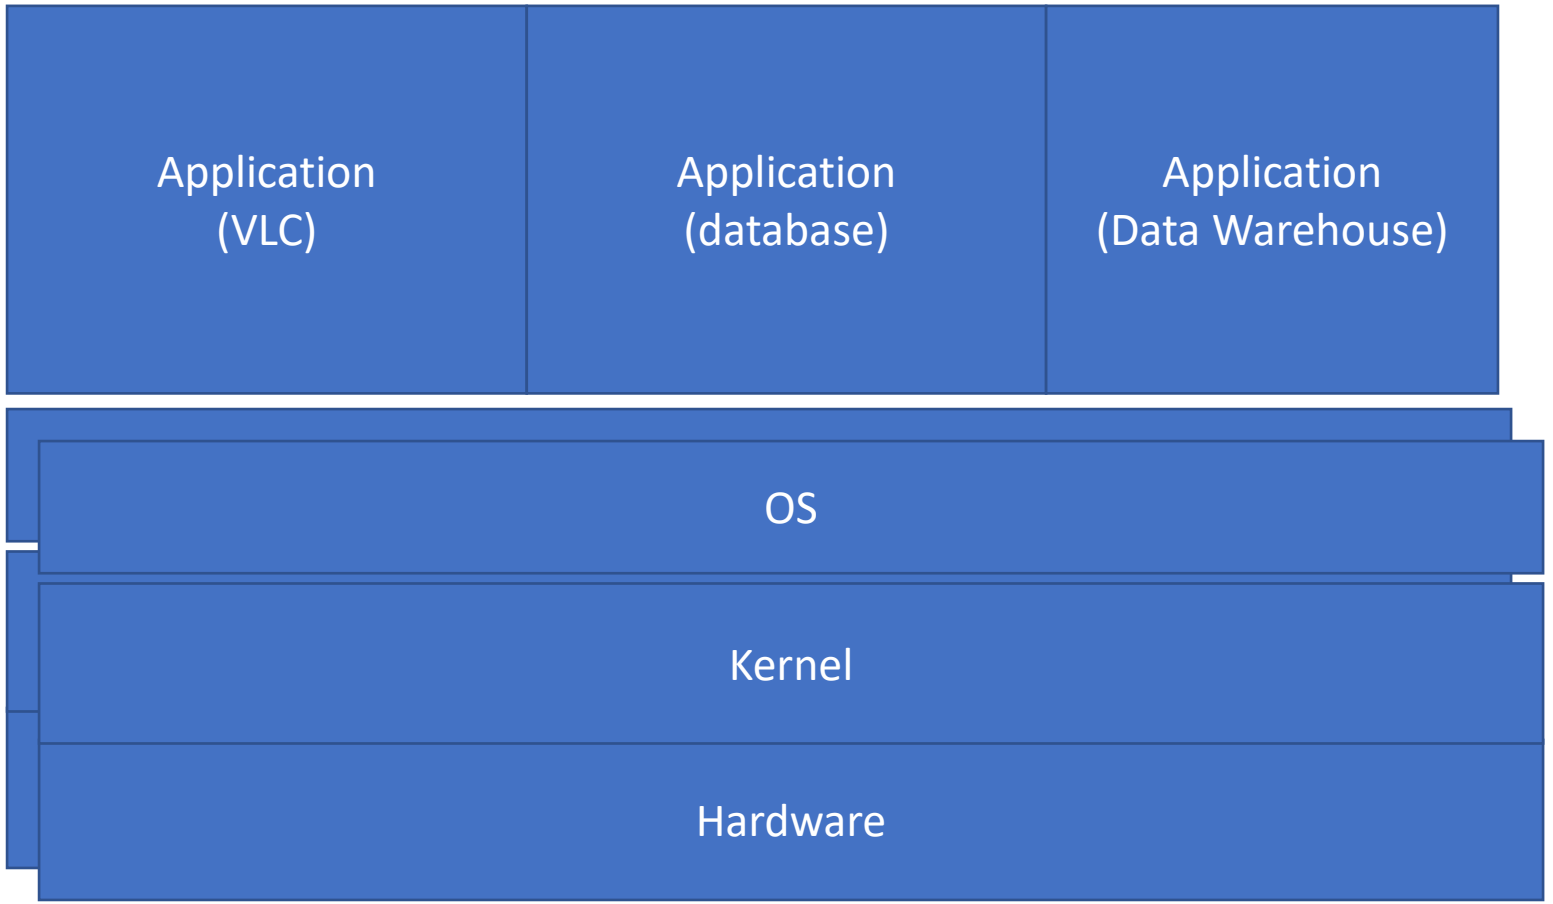
Application
(VLC)
Application
(database)
Application
(Data Warehouse)
OS
Kernel
Hardware

Bigdata -> Concept

- Hadoop
- Spark

Hadoop Frame Work

- Java Based Cluster application ->to deal with big data
- Is an open source tool

History of Hadoop

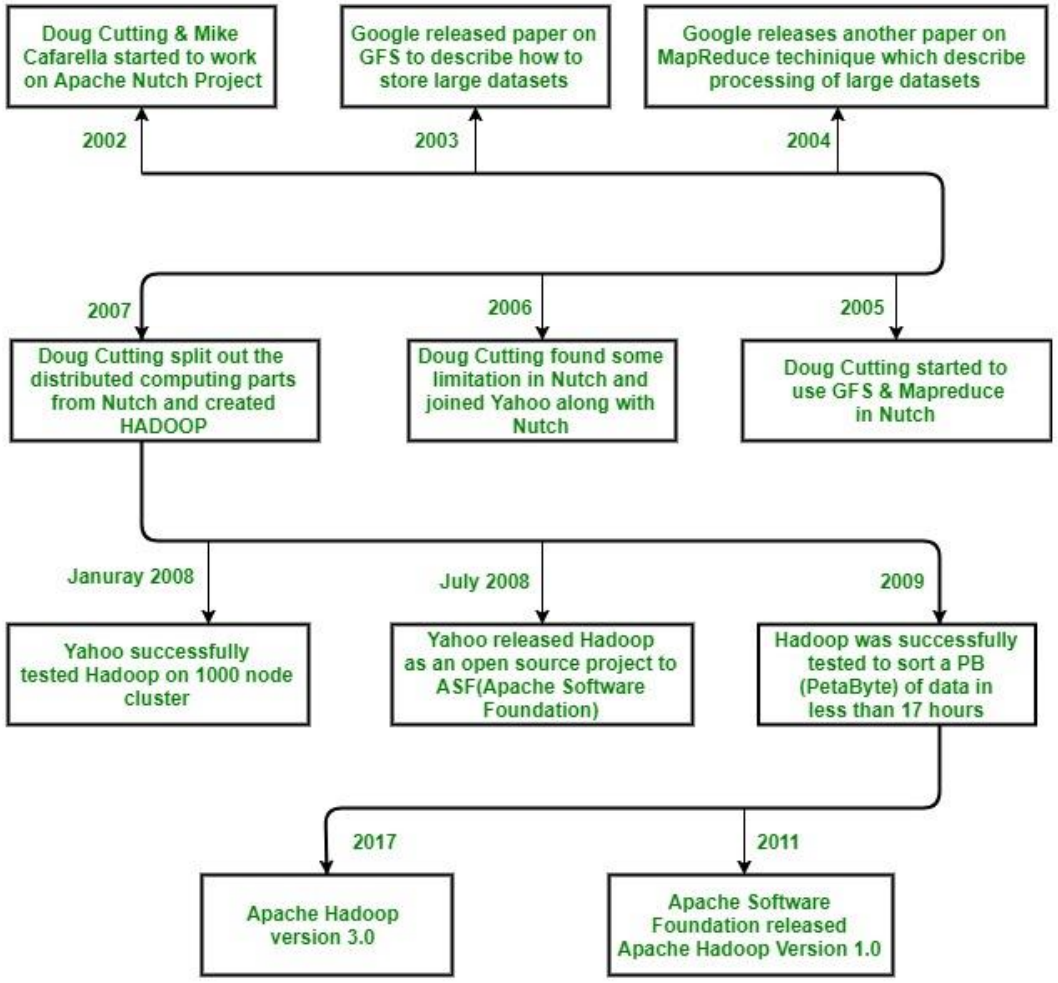

Hadoop Framework

- Distributed Storage(HDFS -> Hadoop Distributed File system)
- Distributed Computing(MapReduce)
- Architecture -> JAVA
- Why JAVA?
- Code -> Compile -> .java -> .class -> JVM(Java Virtual Machine)
- Process -> threads

Hadoop Version 1

- TASK -> Commands -> JAVA Only
- TASK(Data Analytical) -> input(data -> HDFS) -> Command(JAVA -> MAPREDUCE) -> Output(data -> HDFS)
- Storage(HDFS)
  - Master -> Name Node
  - Slave -> Data Node
- CPU+RAM(MapReduce)
  - Master -> Job Tracker
  - Slave -> Task Tacker

Hadoop Version 2

- MapReduce issues
- Communication -> they did the major changes on MapReduce -> YARN(Yet Another Resource Negotiator) -> it will be supported for multiple API and Programming languages
- Way it will handle the job and resource allocation -> Job
- Storage(HDFS)
  - Master -> Name Node
  - Slave -> Data Node
- CPU+RAM(MapReduce)
  - Master -> Resource Manager
  - Slave -> Node Manager

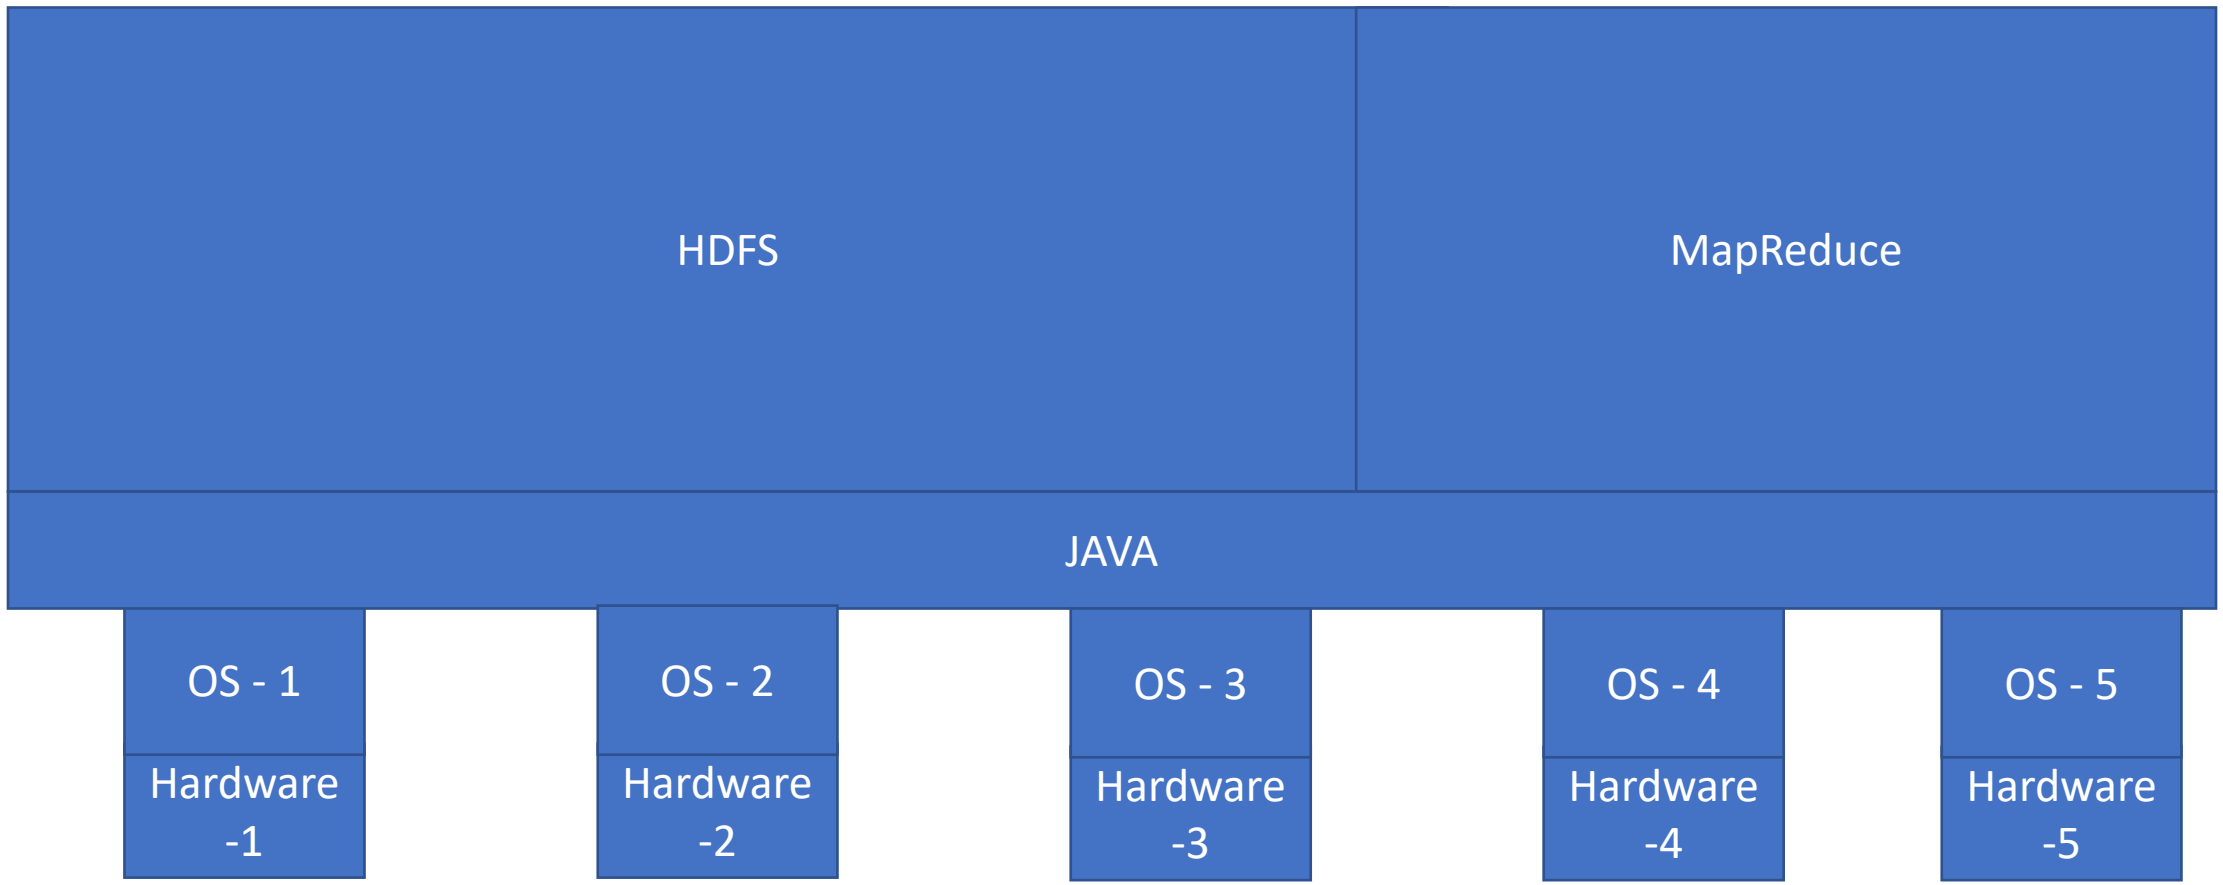
HDFS
MapReduce
JAVA
OS - 1
Hardware -1
OS - 2
Hardware -2
OS - 3
Hardware -3
OS - 4
Hardware -4
OS - 5
Hardware -5

MapReduce V1

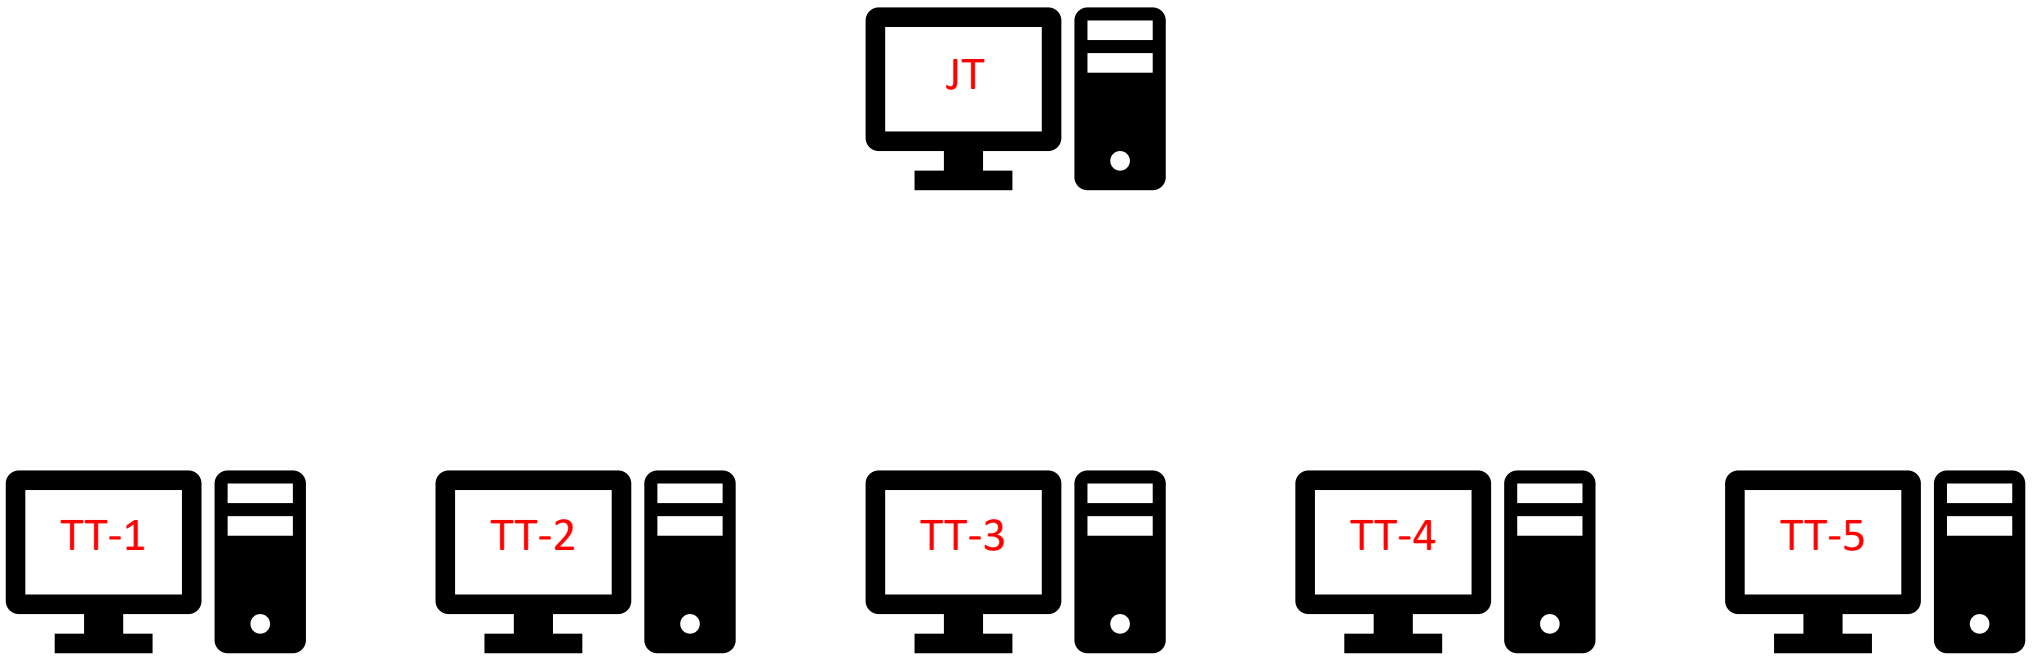

MapReduce V1

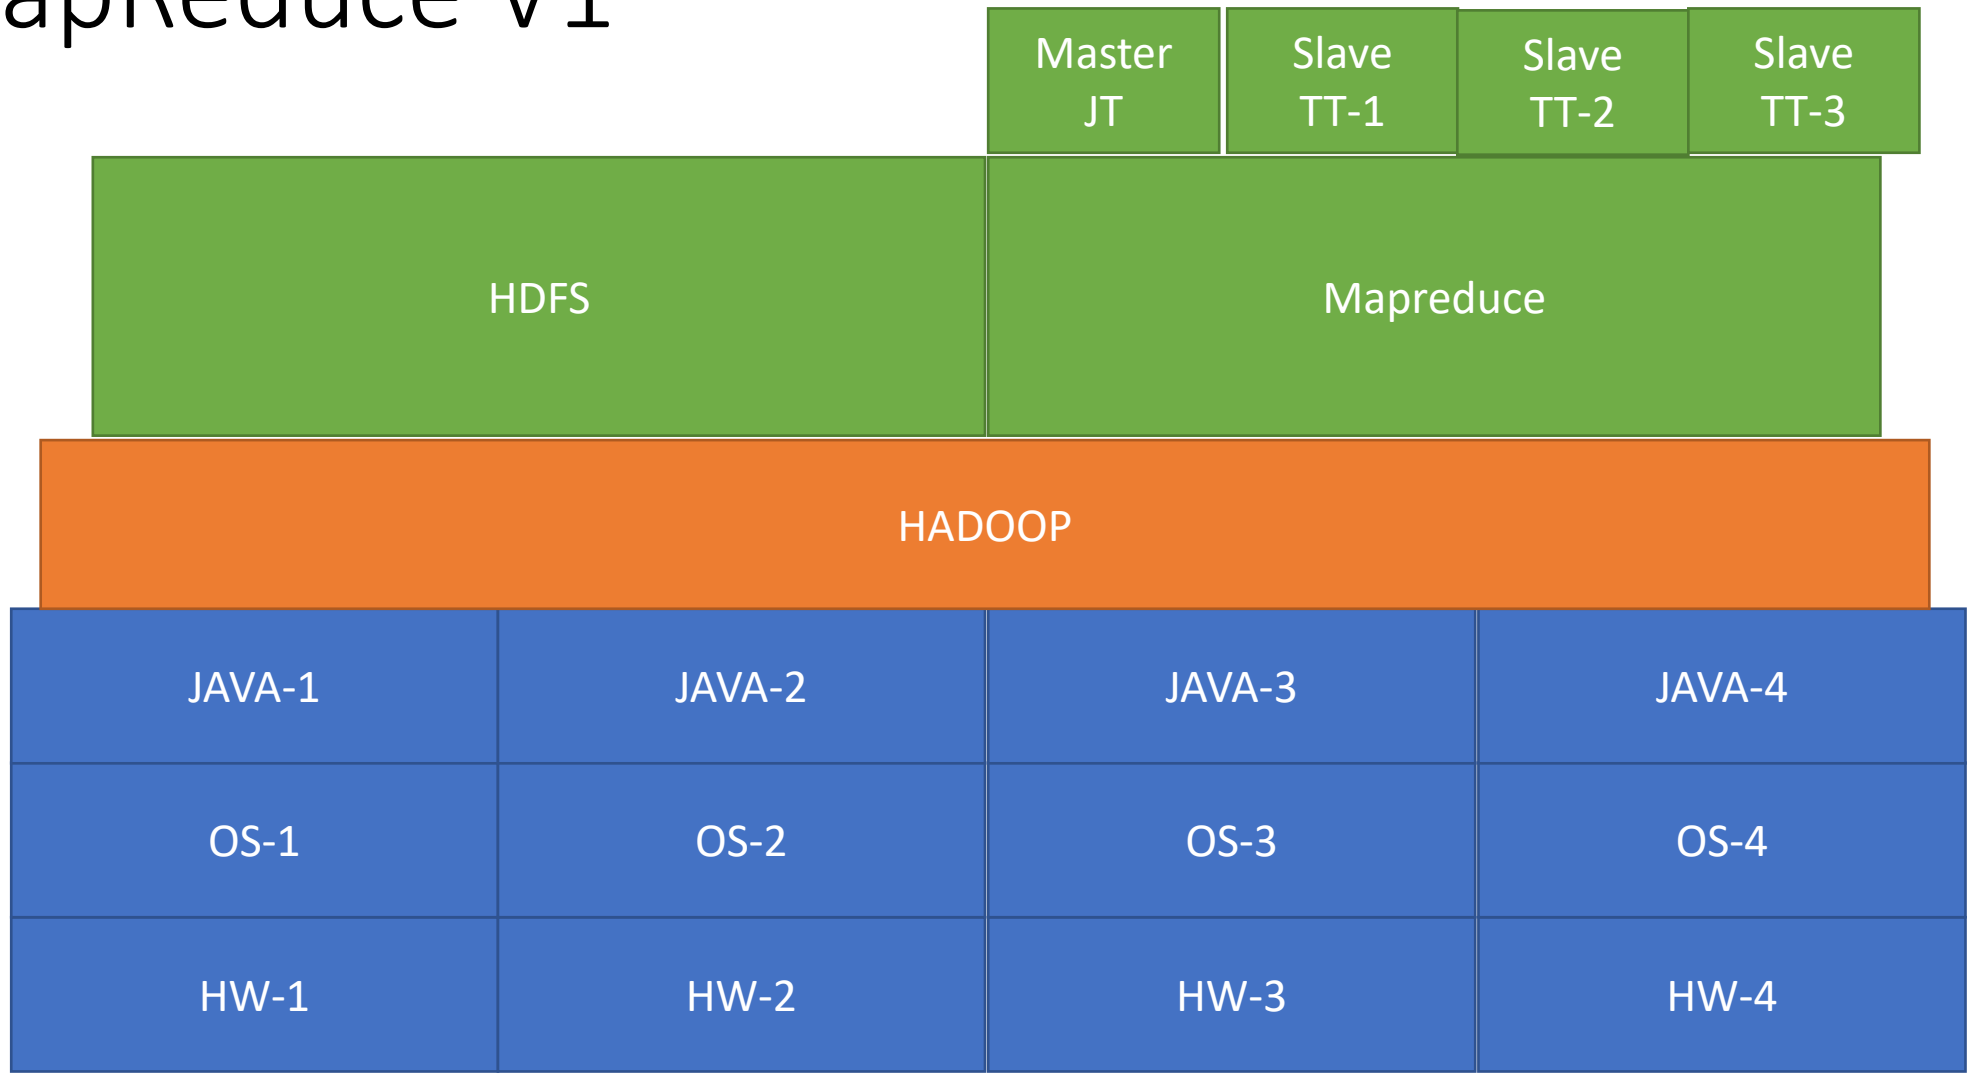

MapReduce V2

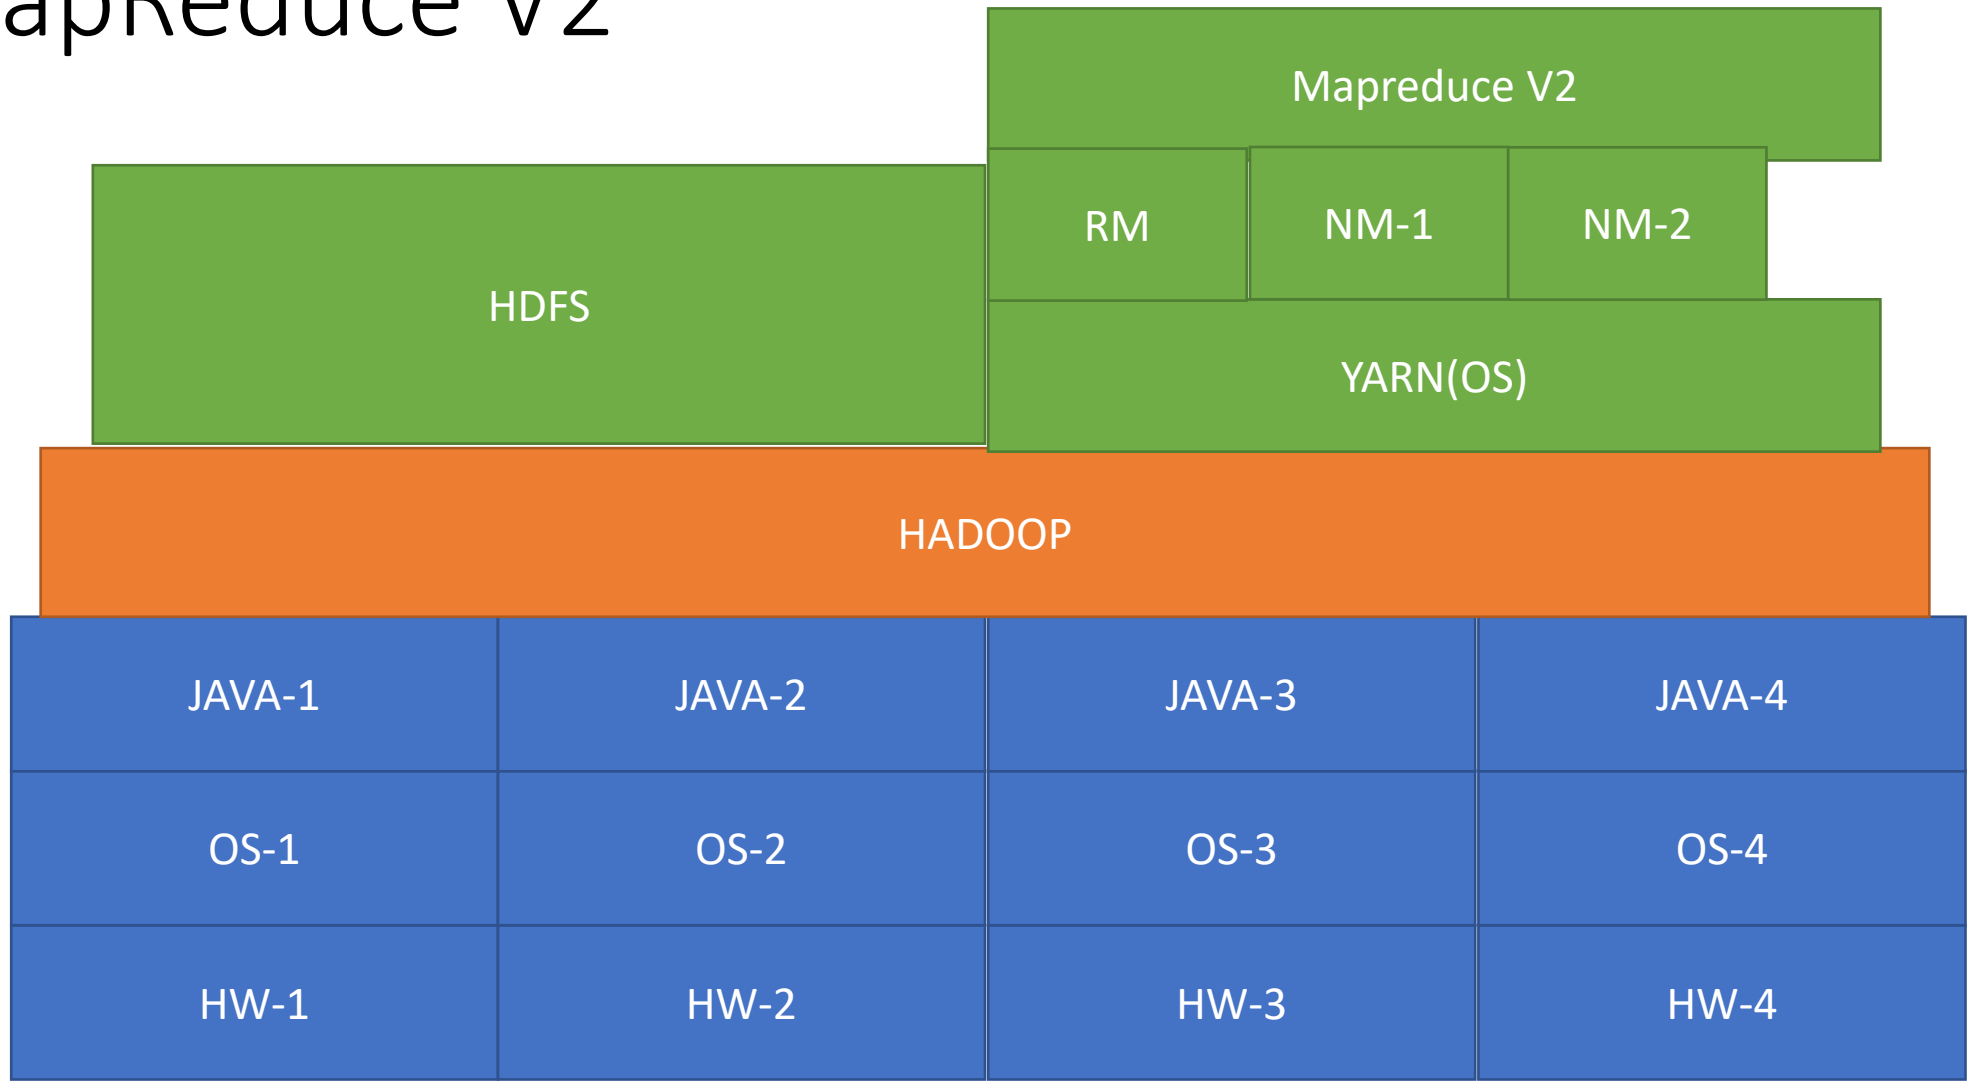

MapReduce V1

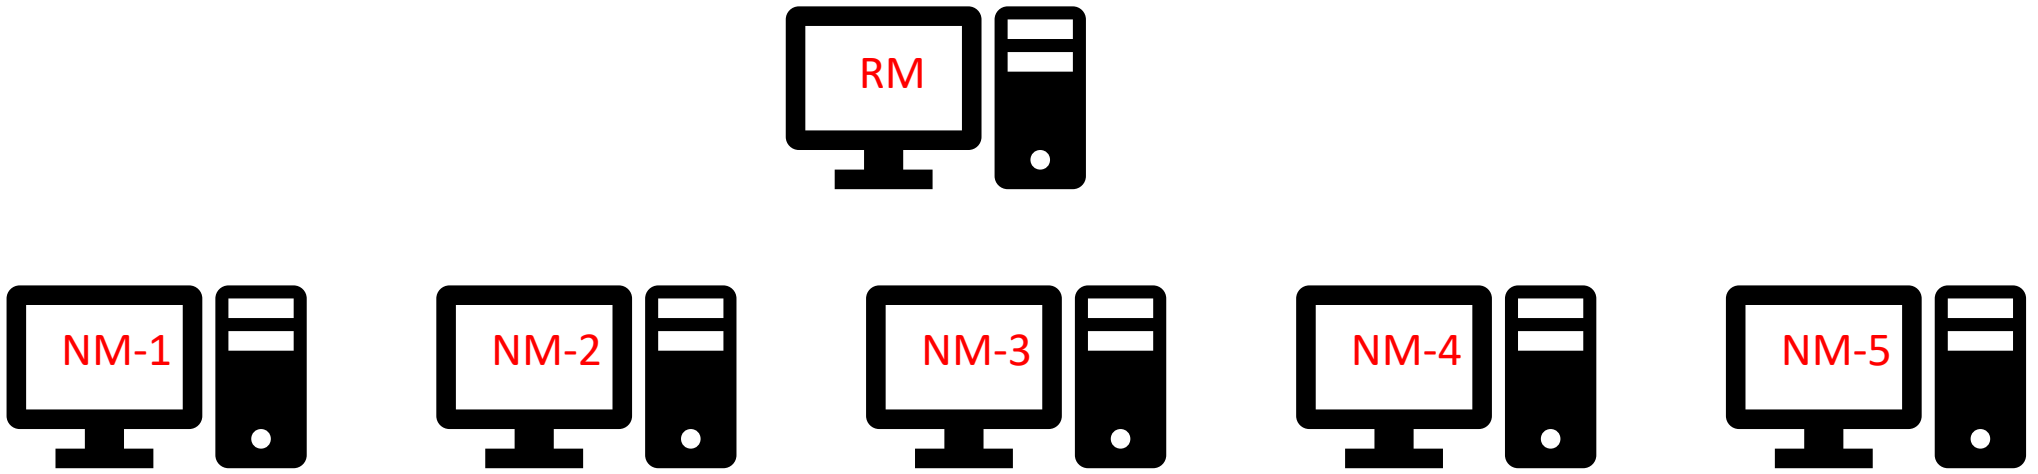

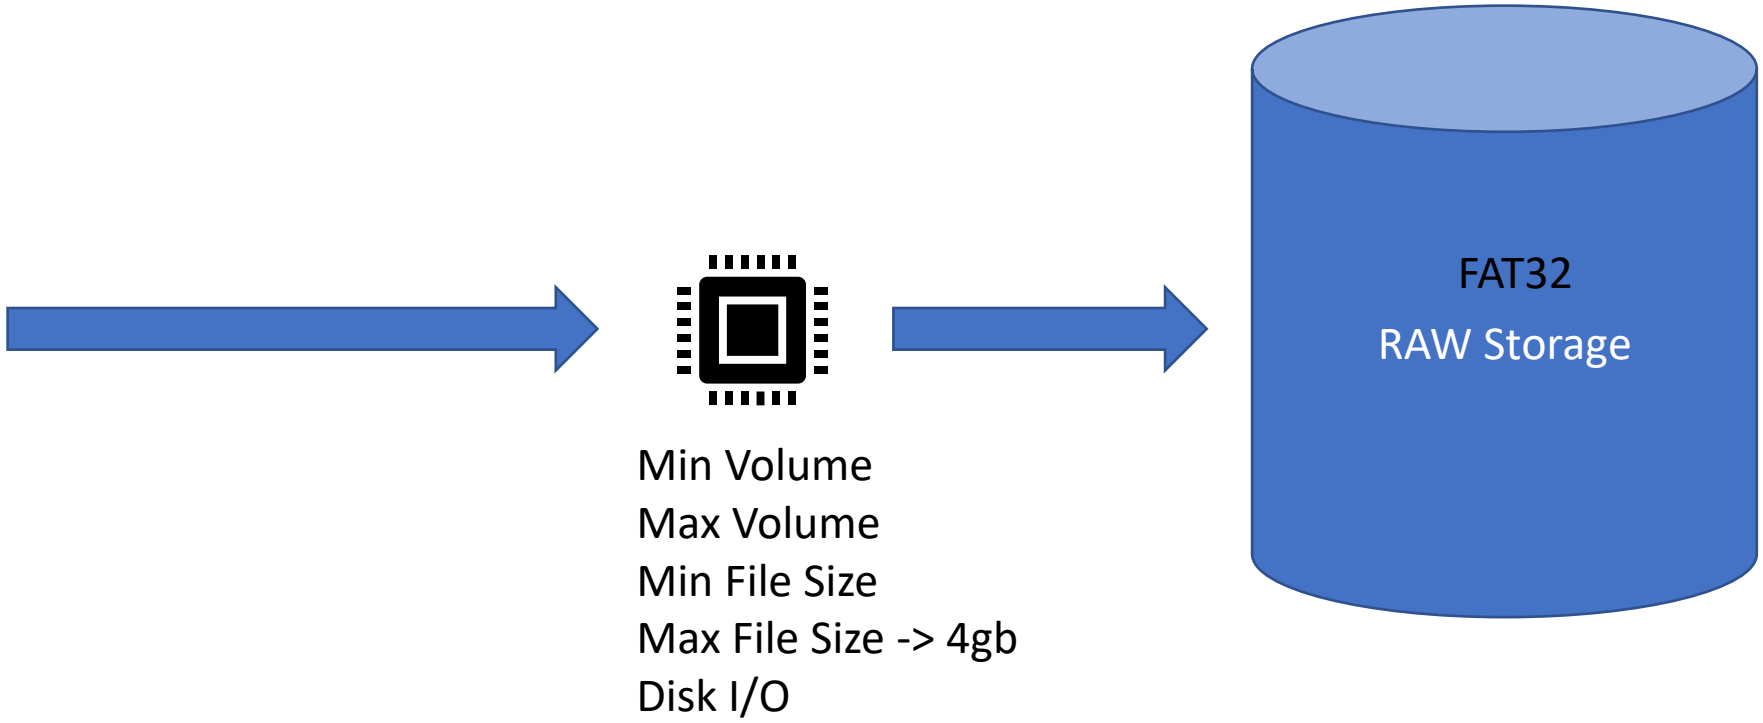
FAT32
RAW Storage
Min Volume
Max Volume
Min File Size
Max File Size -> 4gb
Disk I/O

HDFS

- Distributed File System -> Hadoop Framework -> Resources(Multiple Systems) -> Raw Storage(JAVA)
- If you want to us that raw storage -> Format -> HDFS(JAVA)
- File System -> distributed Storage
- Failover
- No min and max volume, no min and max file size
- HDFS -> Blocks -> Default
  - V1 -> 64MB
  - V2 -> 128MB
- Redundancy -> 1+2 Replication Factor -> Rack Awareness Algorithm
  - V2 -> 1 PB * 2 -> 3 PB
  - V3 -> 6GB -> 9GB

- 500 MB
- Version 1 -> ? Blocks
- Version 2 -> ? Blocks

Functions of Name Node

- It maintains the master daemon that maintains and manages the Data Nodes(Slaves)
- It records the meta data of all the files stored in the cluster e.g the location of the blocks stored, the size of the files, permissions, Hierarchy etc..
  - There are 2 special files are associated with the meta data
  - FsImage :- it contains the complete state of the file system namespace since the start of the Name node
  - EditLogs : - it contains all the recent modifications made to the file system with respect to the most recent FsImage

Functions of Name Node

- It records each change that take place to the system metadata ex: if the file is deleted in HDFS, the name node will immediately record in the EditLogs
- It regularly receives a Heartbeat and a block report form all the data nodes in the cluster to ensure the data node is alive
- It keeps a record of all the blocks in HDFS and which nodes these bocks are located
- The name node is also responsible to take care of the replication factor of all the blocks
- In case of the data node failure the name nodes choses new data node for new replicas, balance the disk usage and manages the communication traffic to the data nodes

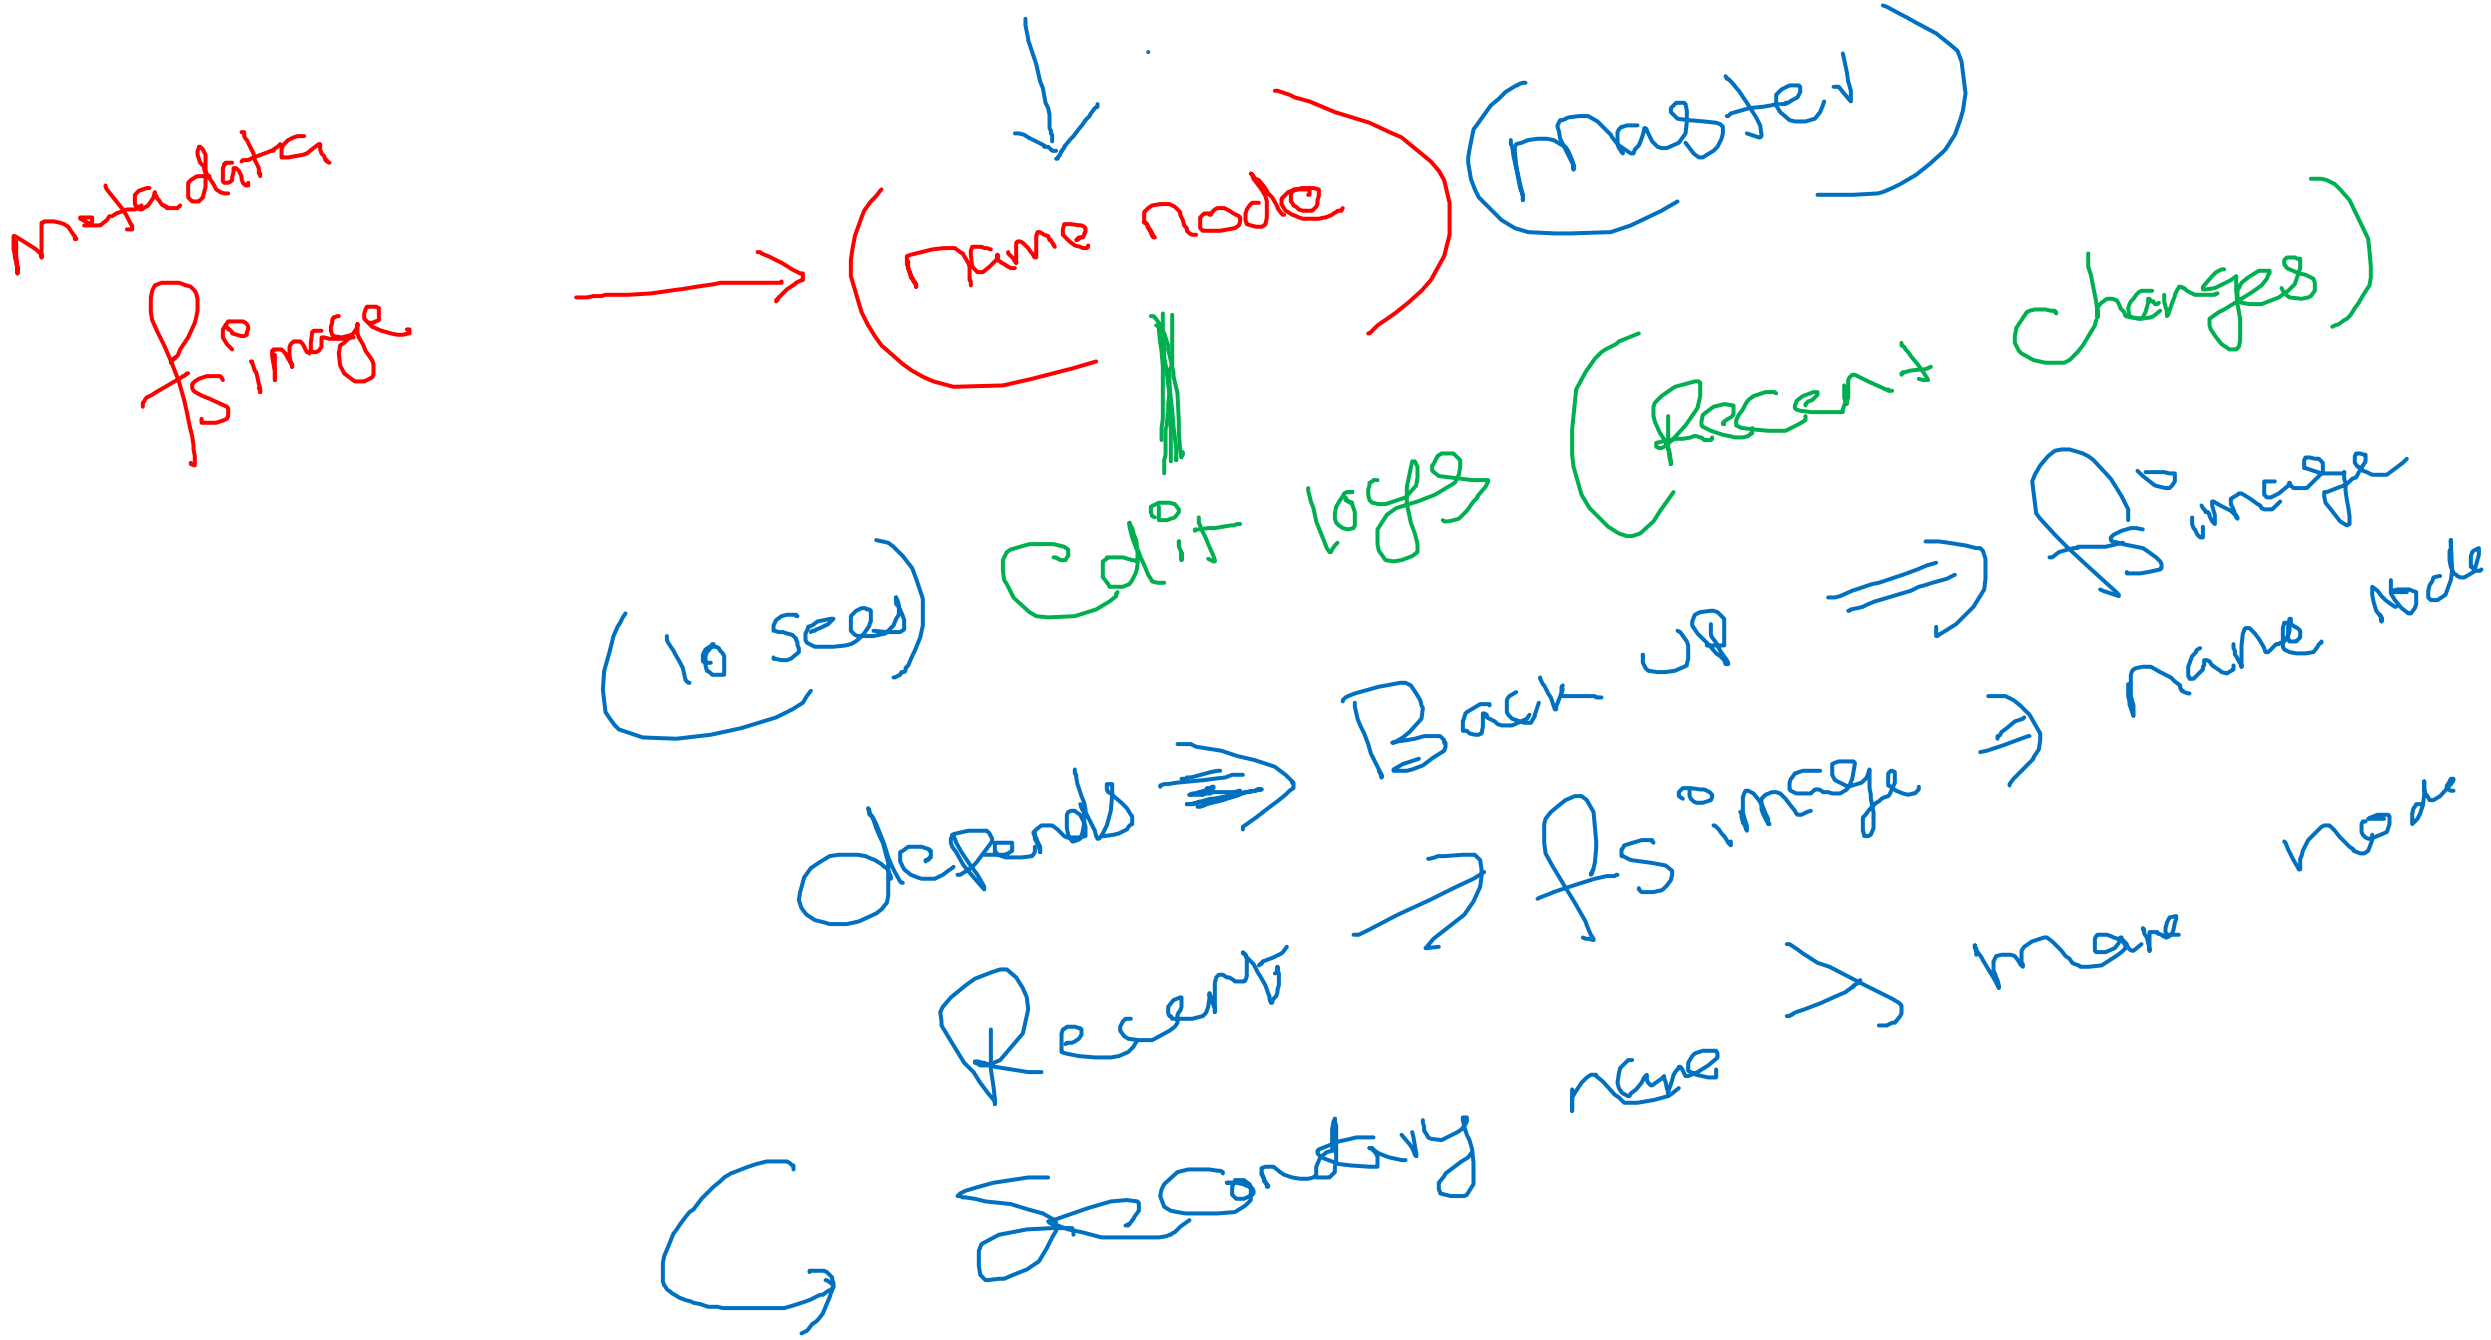
Metadata
fsimage
name node
(master)
edit logs
(recent changes)
(to see)
depends
Back up
fsimage
name node
Recent
fsimage
name node
Secondary name node

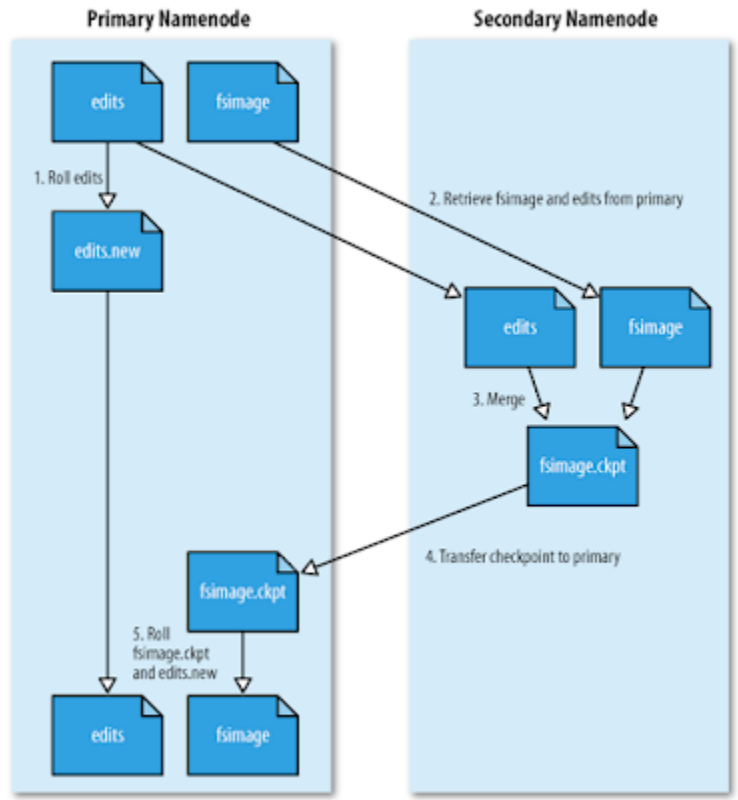
Primary Namenode
Secondary Namenode
edits
fsimage
1. Roll edits
edits.new
2. Retrieve fsimage and edits from primary
edits
fsimage
3. Merge
fsimage.ckpt
4. Transfer checkpoint to primary
fsimage.ckpt
5. Roll
fsimage.ckpt
and edits.new
edits
fsimage

Functions of Data Node

- These are slave daemons which runs on each slave machine
- The actual data is stored on data nodes
- The data nodes perform the low-level read and write requests from the clients directly
- They send the heart beats to the name node periodically to report overall health of HDFS, by default this frequency is set to 3 seconds

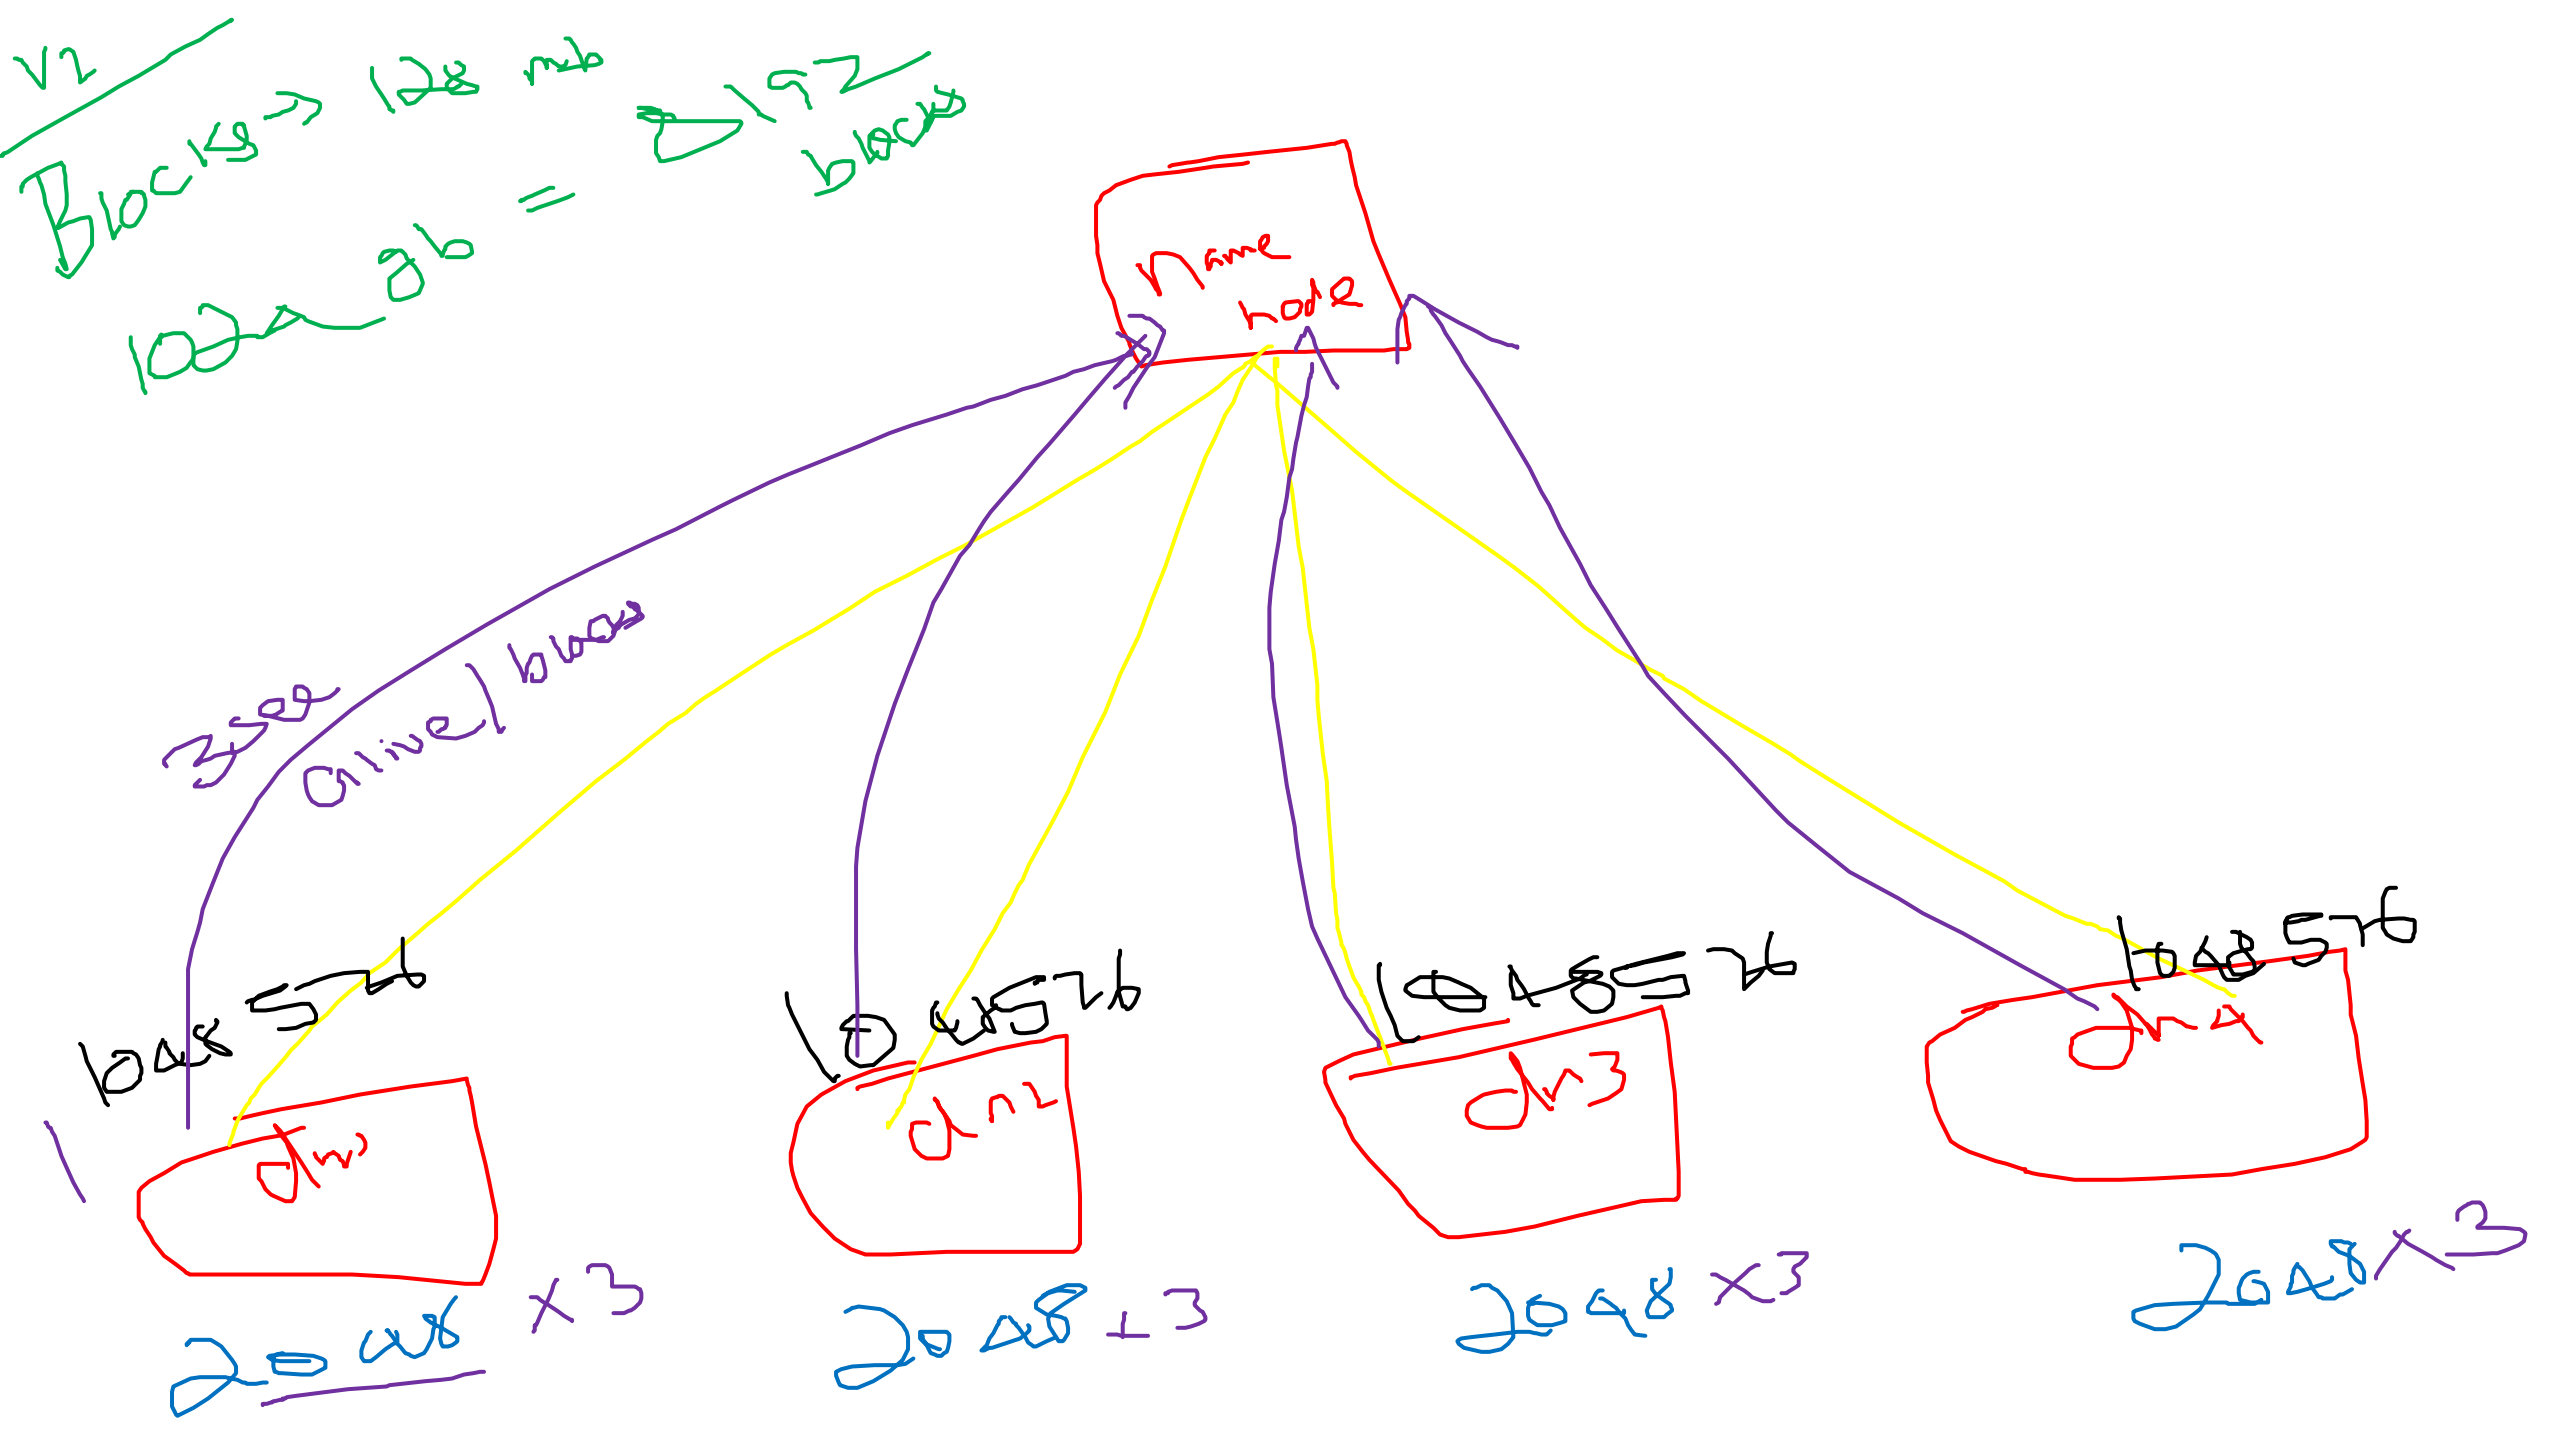
V2
Blocks -> 128 mb
1024 gb = 8192 blocks
Name node
3sec alive/ block
1048576
1048576
1048576
1048576
dn1
dn2
dn3
dn4
2048 x3
2048 x3
2048 x3
2048 x3

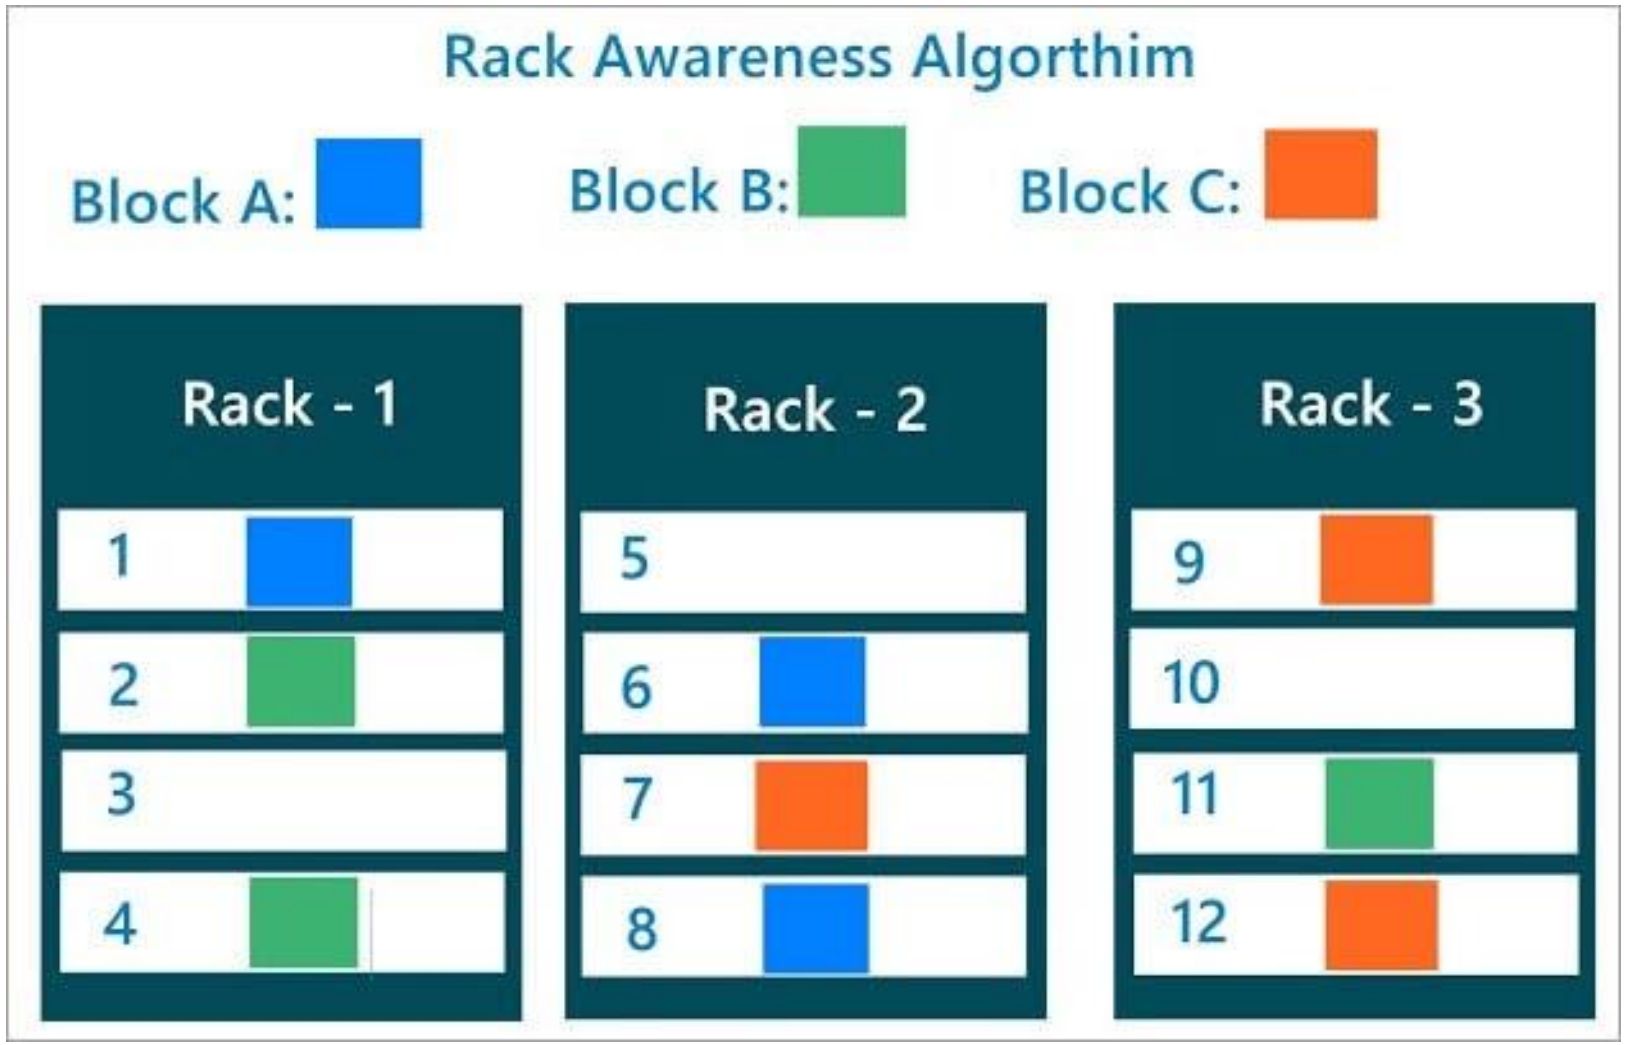
Rack Awareness Algorthim
Block A:
Block B:
Block C:
Rack - 1
1
2
3
4
Rack - 2
5
6
7
8
Rack - 3
9
10
11
12

HDFS -> Read and Write

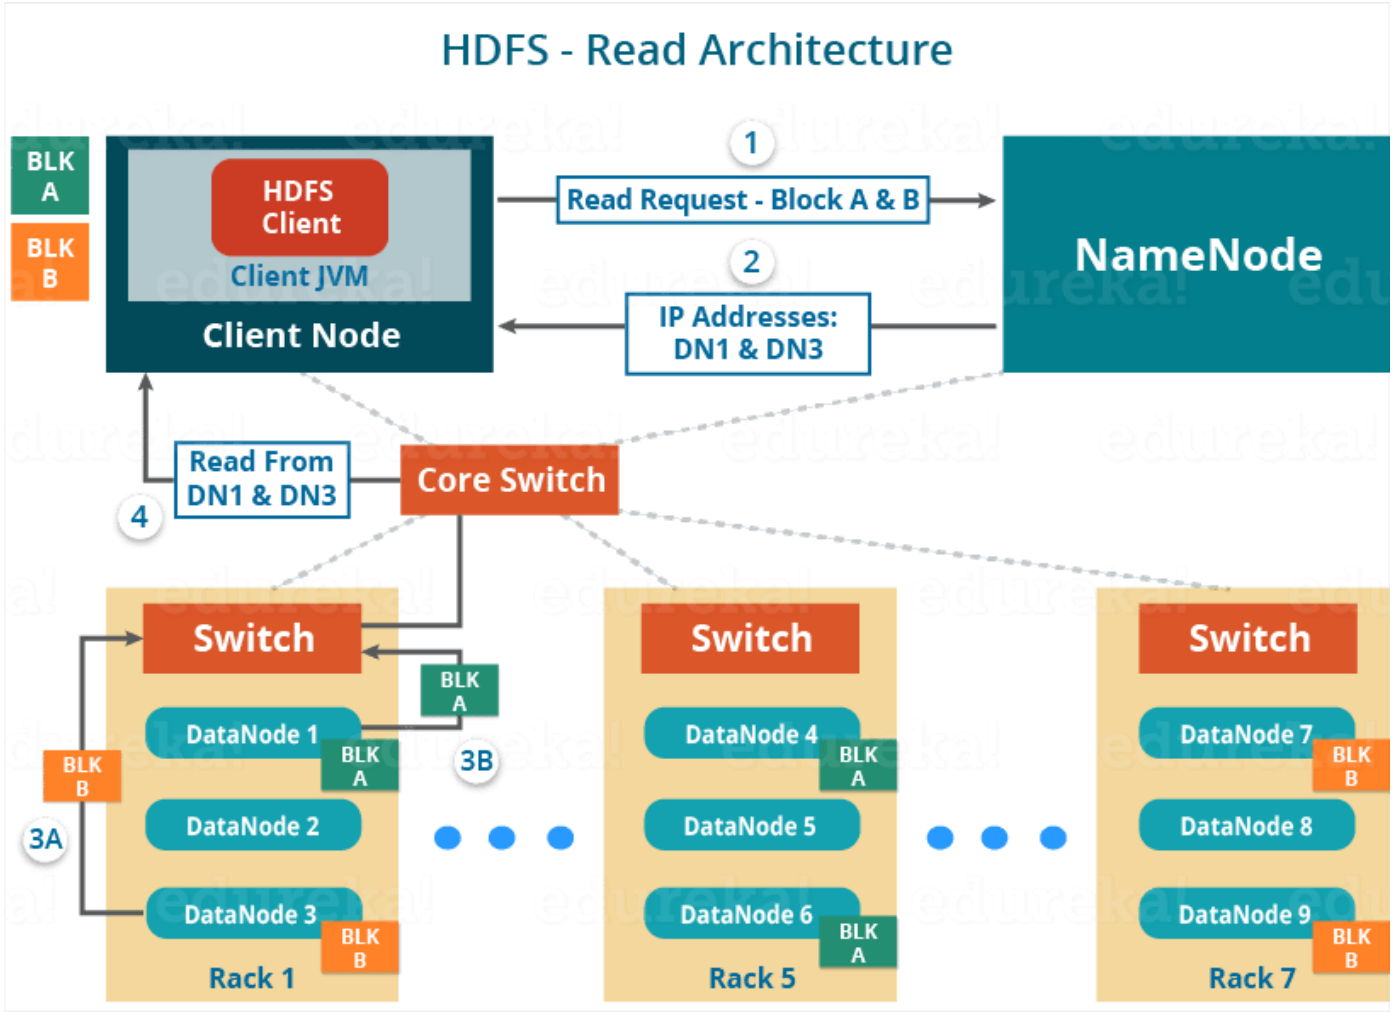

HDFS Write

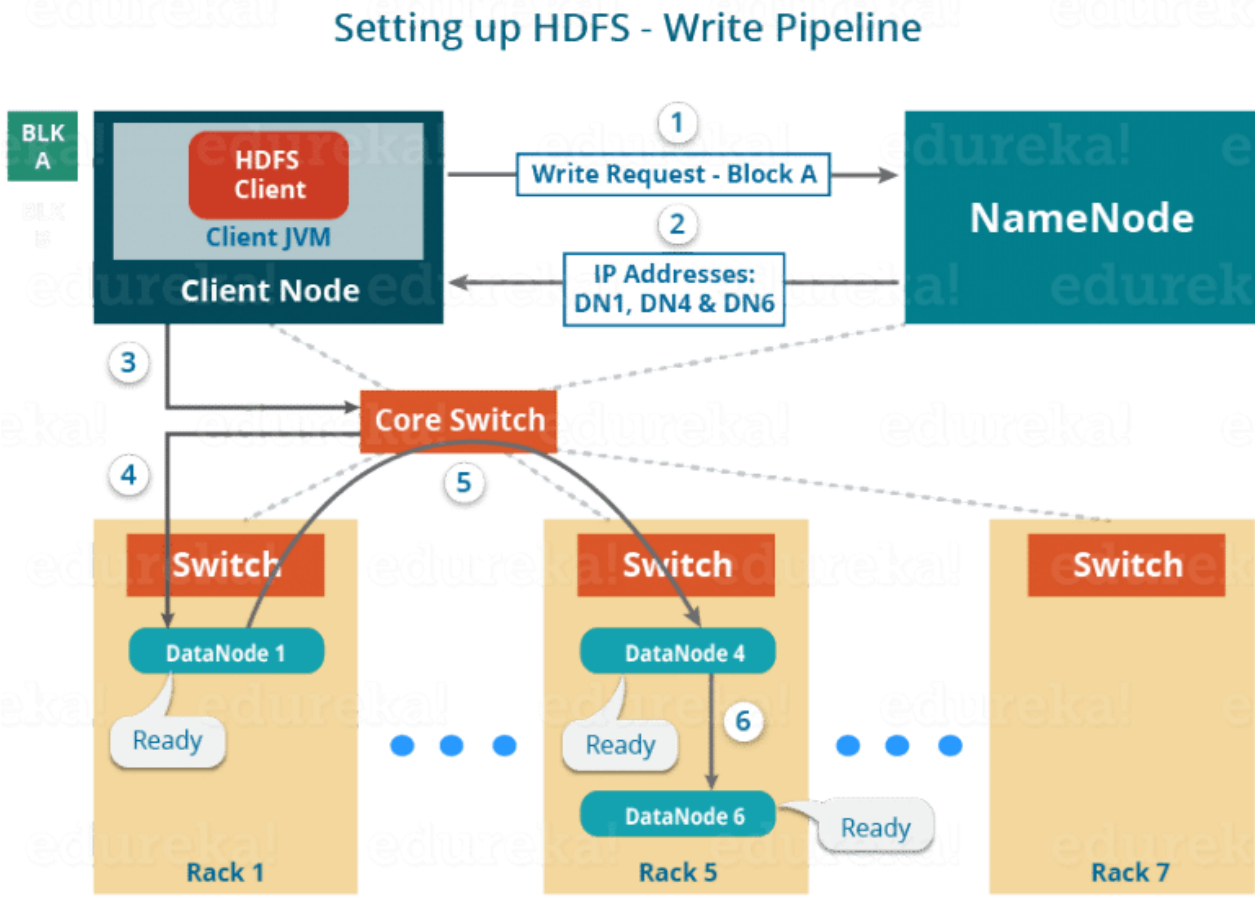
Setting up HDFS - Write Pipeline
BLK A
HDFS Client
Client JVM
Client Node
1
Write Request - Block A
2
IP Addresses: DN1, DN4 & DN6
NameNode
3
Core Switch
4
5
Switch
DataNode 1
Ready
Rack 1
Switch
DataNode 4
Ready
6
DataNode 6
Ready
Rack 5
Switch
Rack 7

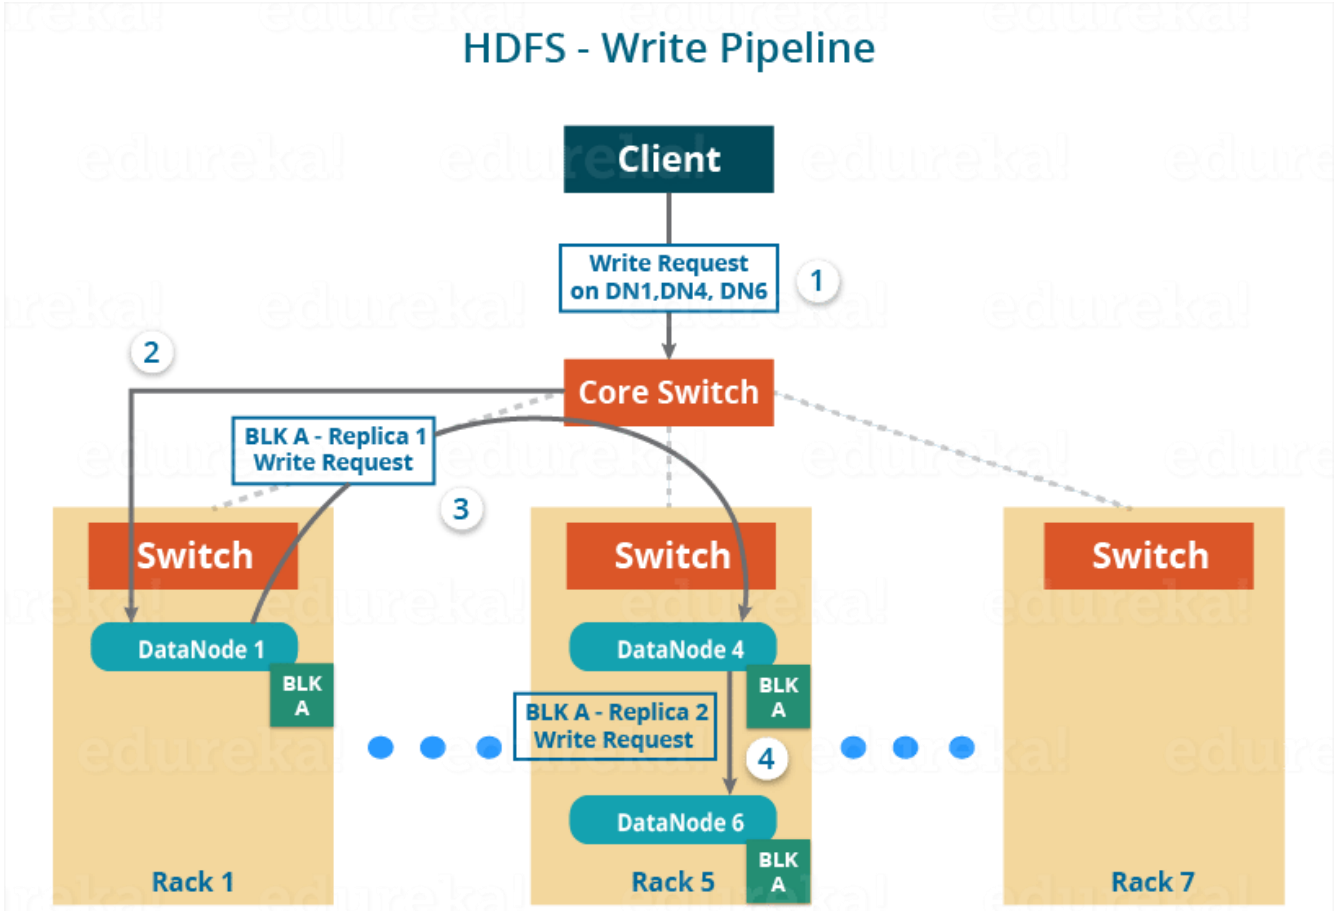
HDFS - Write Pipeline
Client
Write Request
on DN1,DN4, DN6
1
2
Core Switch
BLK A - Replica 1
Write Request
3
Switch
Switch
Switch
DataNode 1
BLK
A
DataNode 4
BLK
A
BLK A - Replica 2
Write Request
4
DataNode 6
BLK
A
Rack 1
Rack 5
Rack 7

- OLAP -> Editing is not allowed

HDFS

MapReduce

- Distributed Computing Component that is available on Hadoop
- If you schedule any Mapreduce task
- V1
  - Job Tracker
  - Task Tracker
- V2
  - Resource Manager
  - Node Manager

How MapReduce Makes it possible?

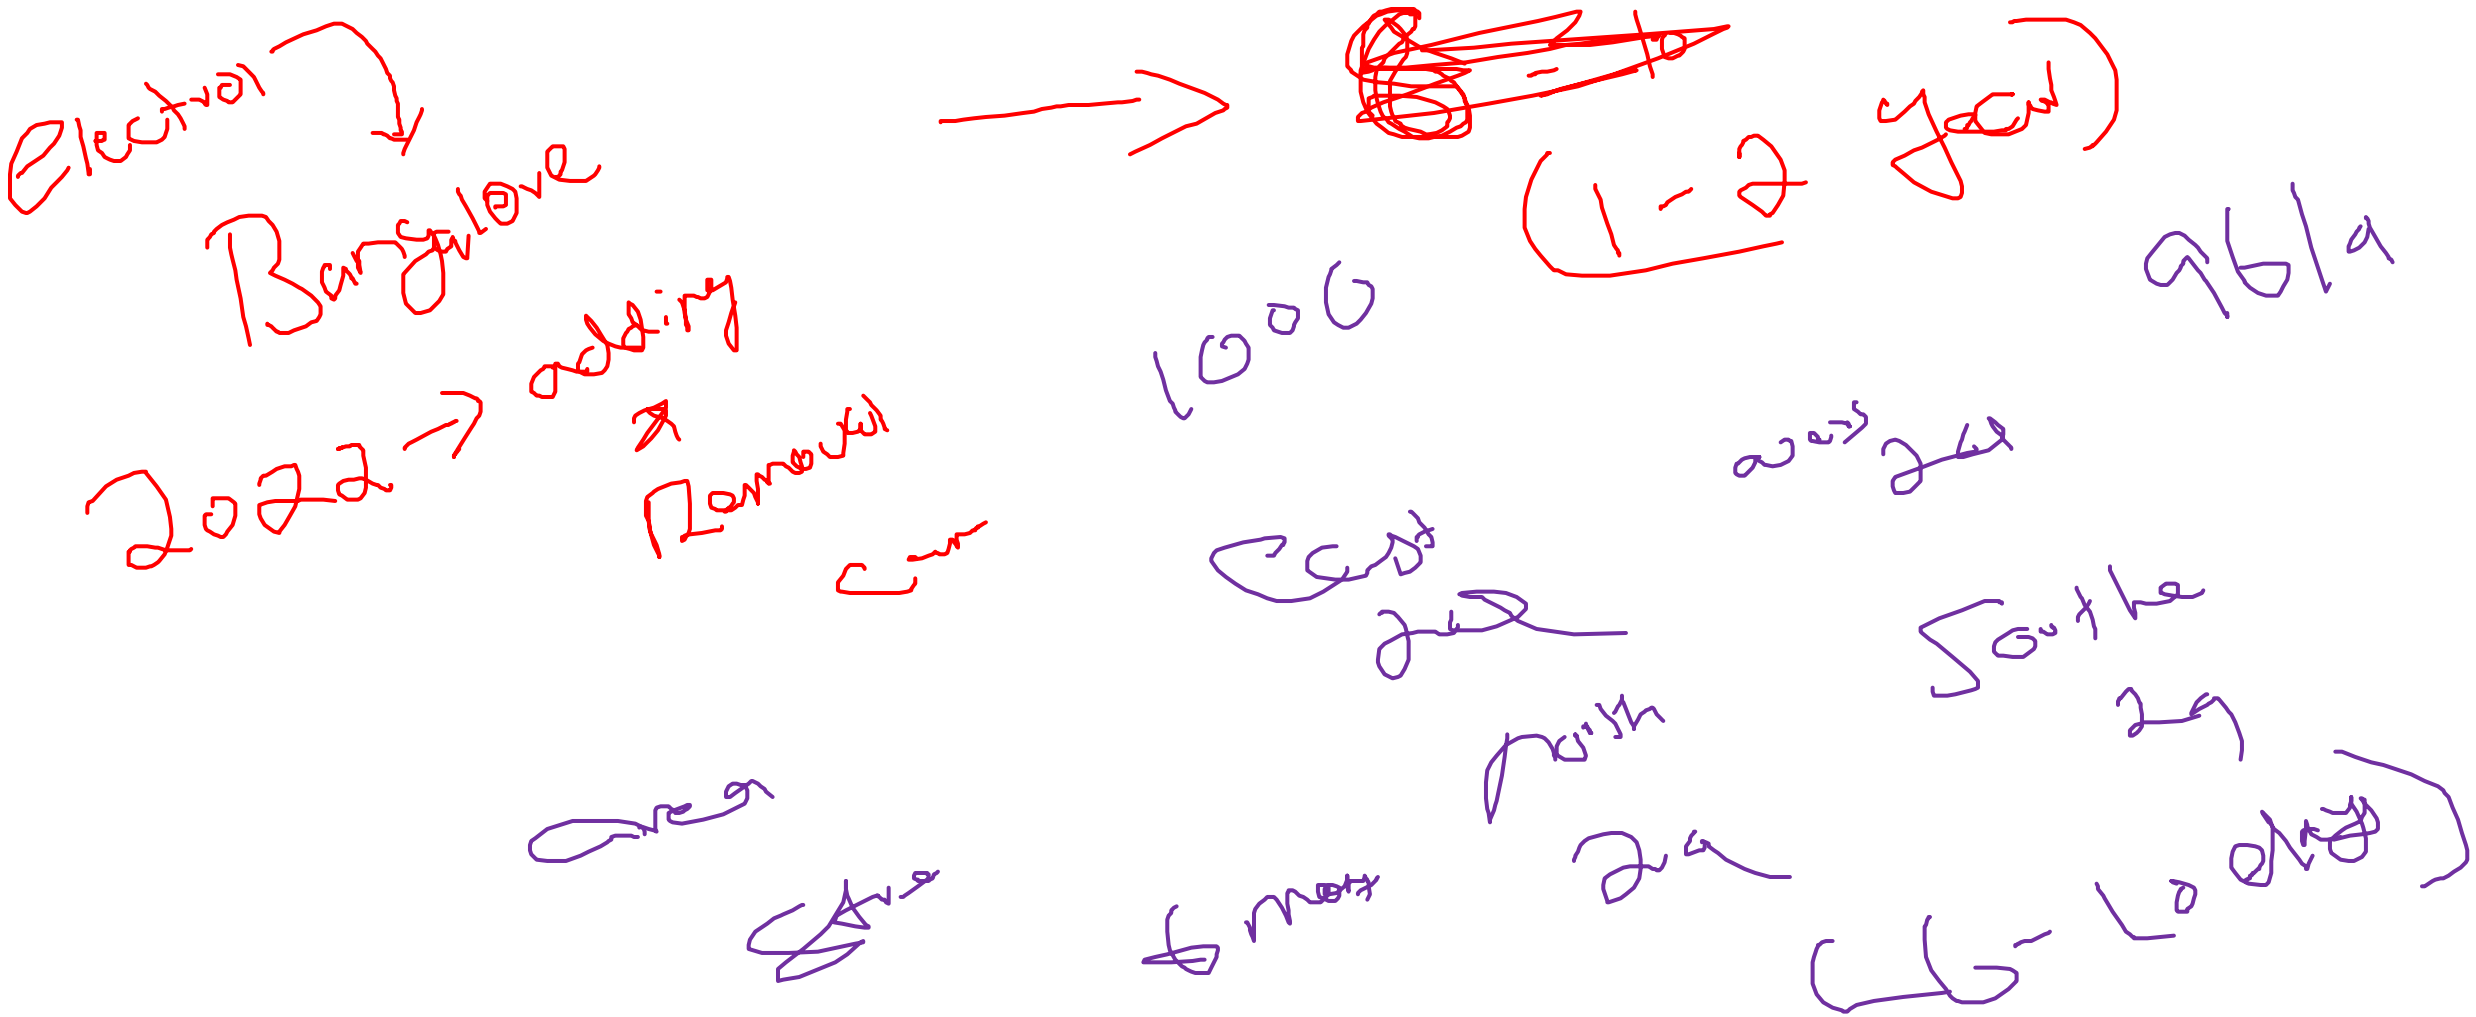

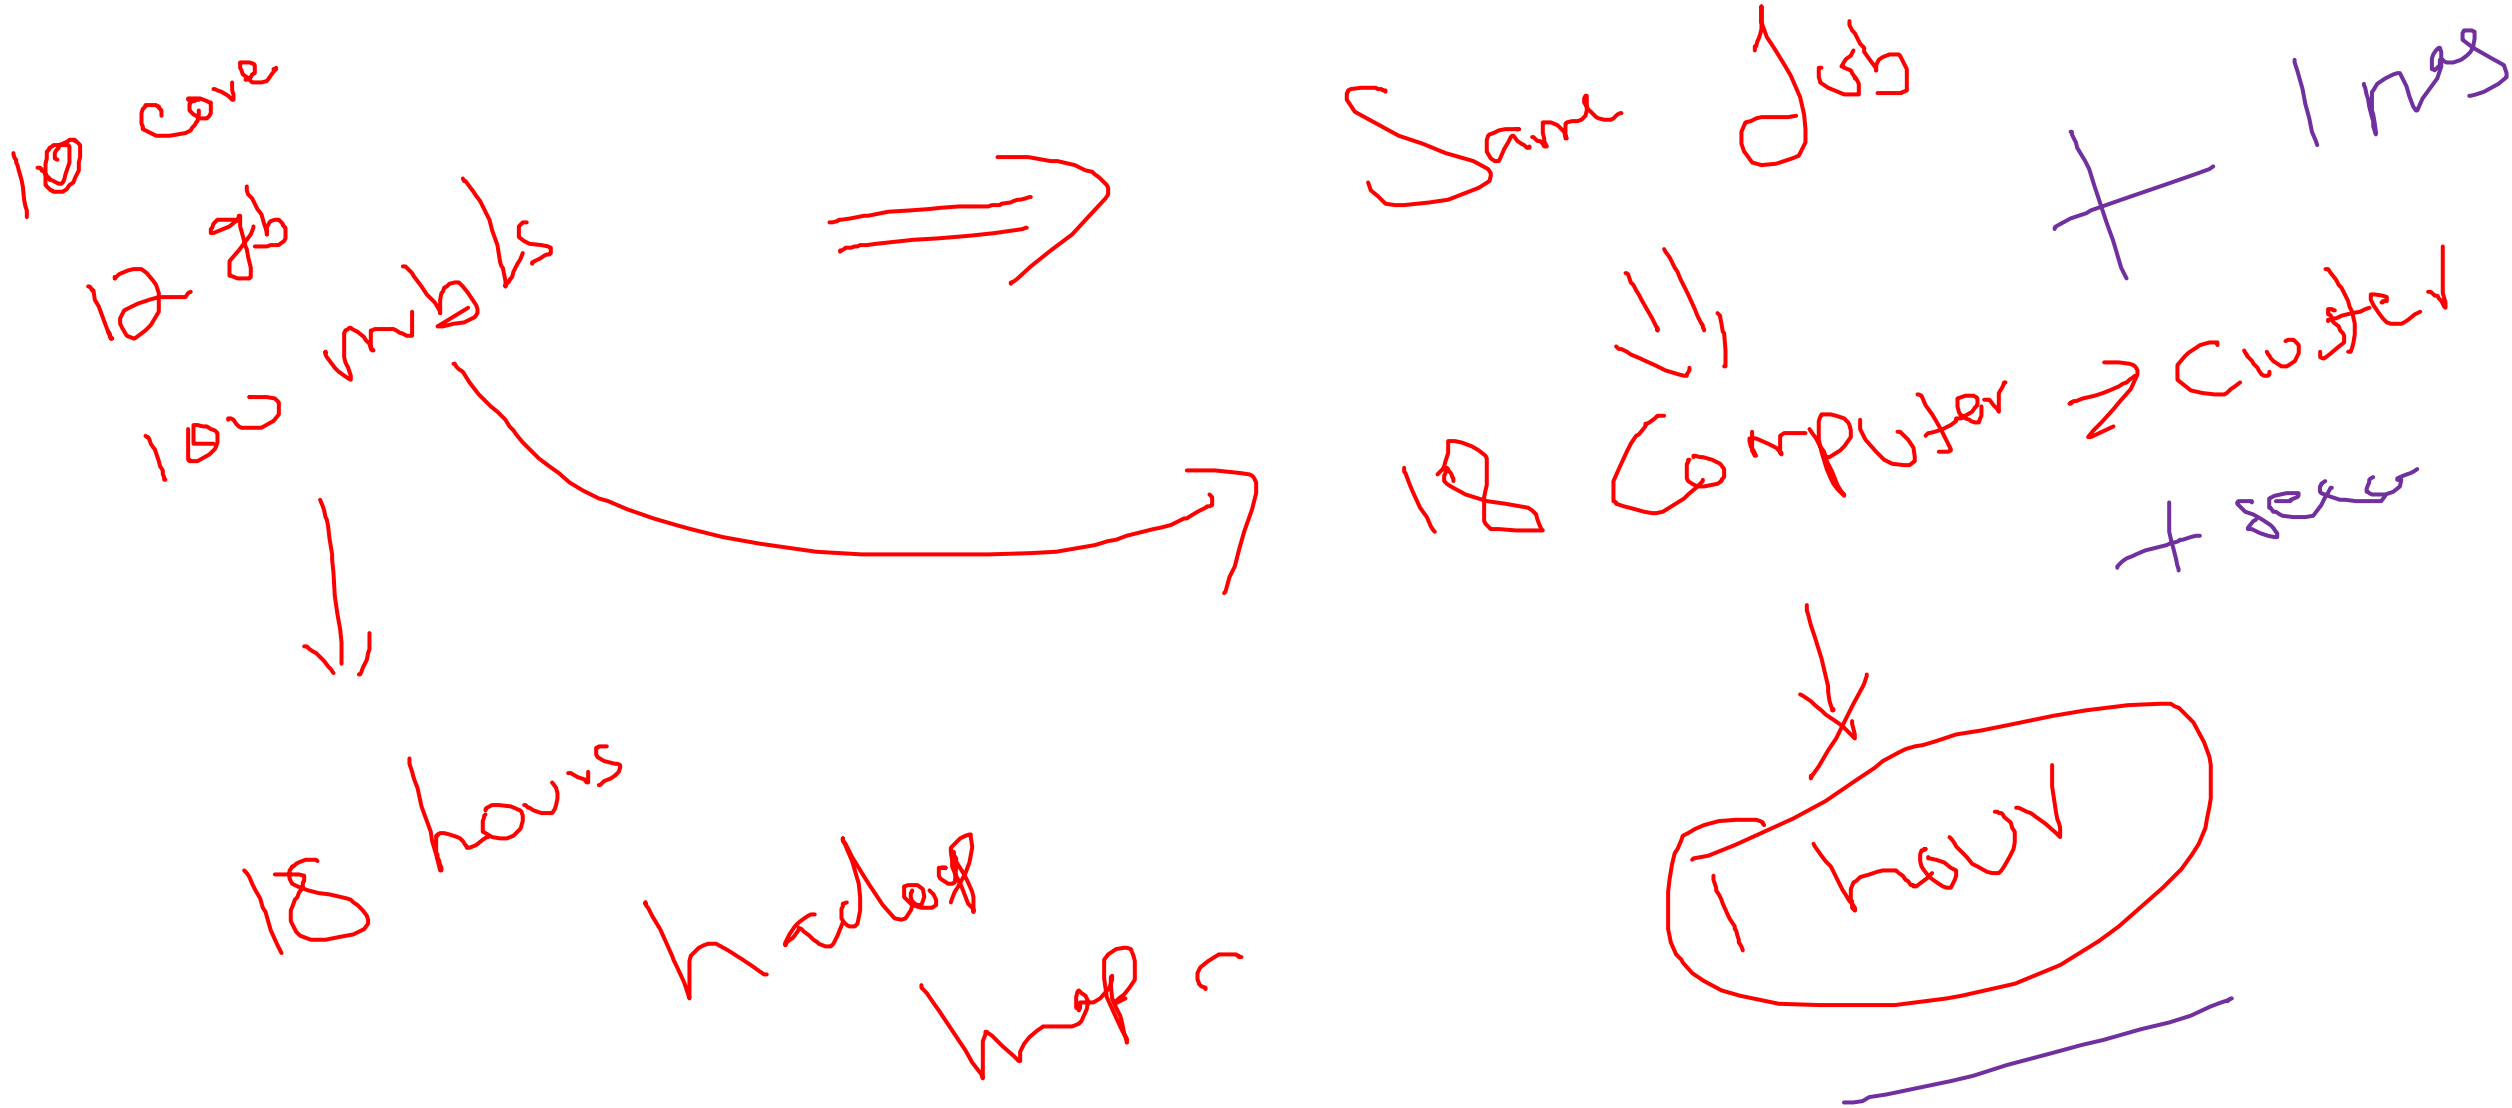
10 core
12 gb
100 mb/s
18 hours
Same Job
10 Computer
1 hour
1 mins
+ setup

YARN(Yet another Resource Negotiator)

- Version 1 -> Mapreduce V1 -> Resource allocation and Job Handling
- In version2 -> Hadoop -> they separated the job handler and the resource allocation by means of 2 different principles
- MapReduce v2 -> Actual job will run on MapReduce v2 under the guidance YARN
- YARN -> Allocating the resources that is required for running the job

YARN(Yet another Resource Negotiator)

- It was developed by the company Horton Works
- Idea was introduced this technology which will act as OS
- YARN -> Distributed OS -> interface to manage the cluster hardware(entire cluster)
- Cluster Capacity -> only from the Worker Nodes

YARN

- Distributed OS
  - Managing the Cluster Hardware(Cluster Capacity)
  - Resource allocation for Apps/Tasks
  - Monitor the health of entire cluster(Worker Node down, it will try to run the job somewhere else, wait for heartbeat from the worker node which is down)
    - Heartbeat -> kind of communication message that is send from Worker node to Master node(inform the worker node is alive)
  - Logging -> log file aggregation property is there
  - Pluggable -> install(Spark, Tez, MapReduce)
  - Prioritization -> priority to job

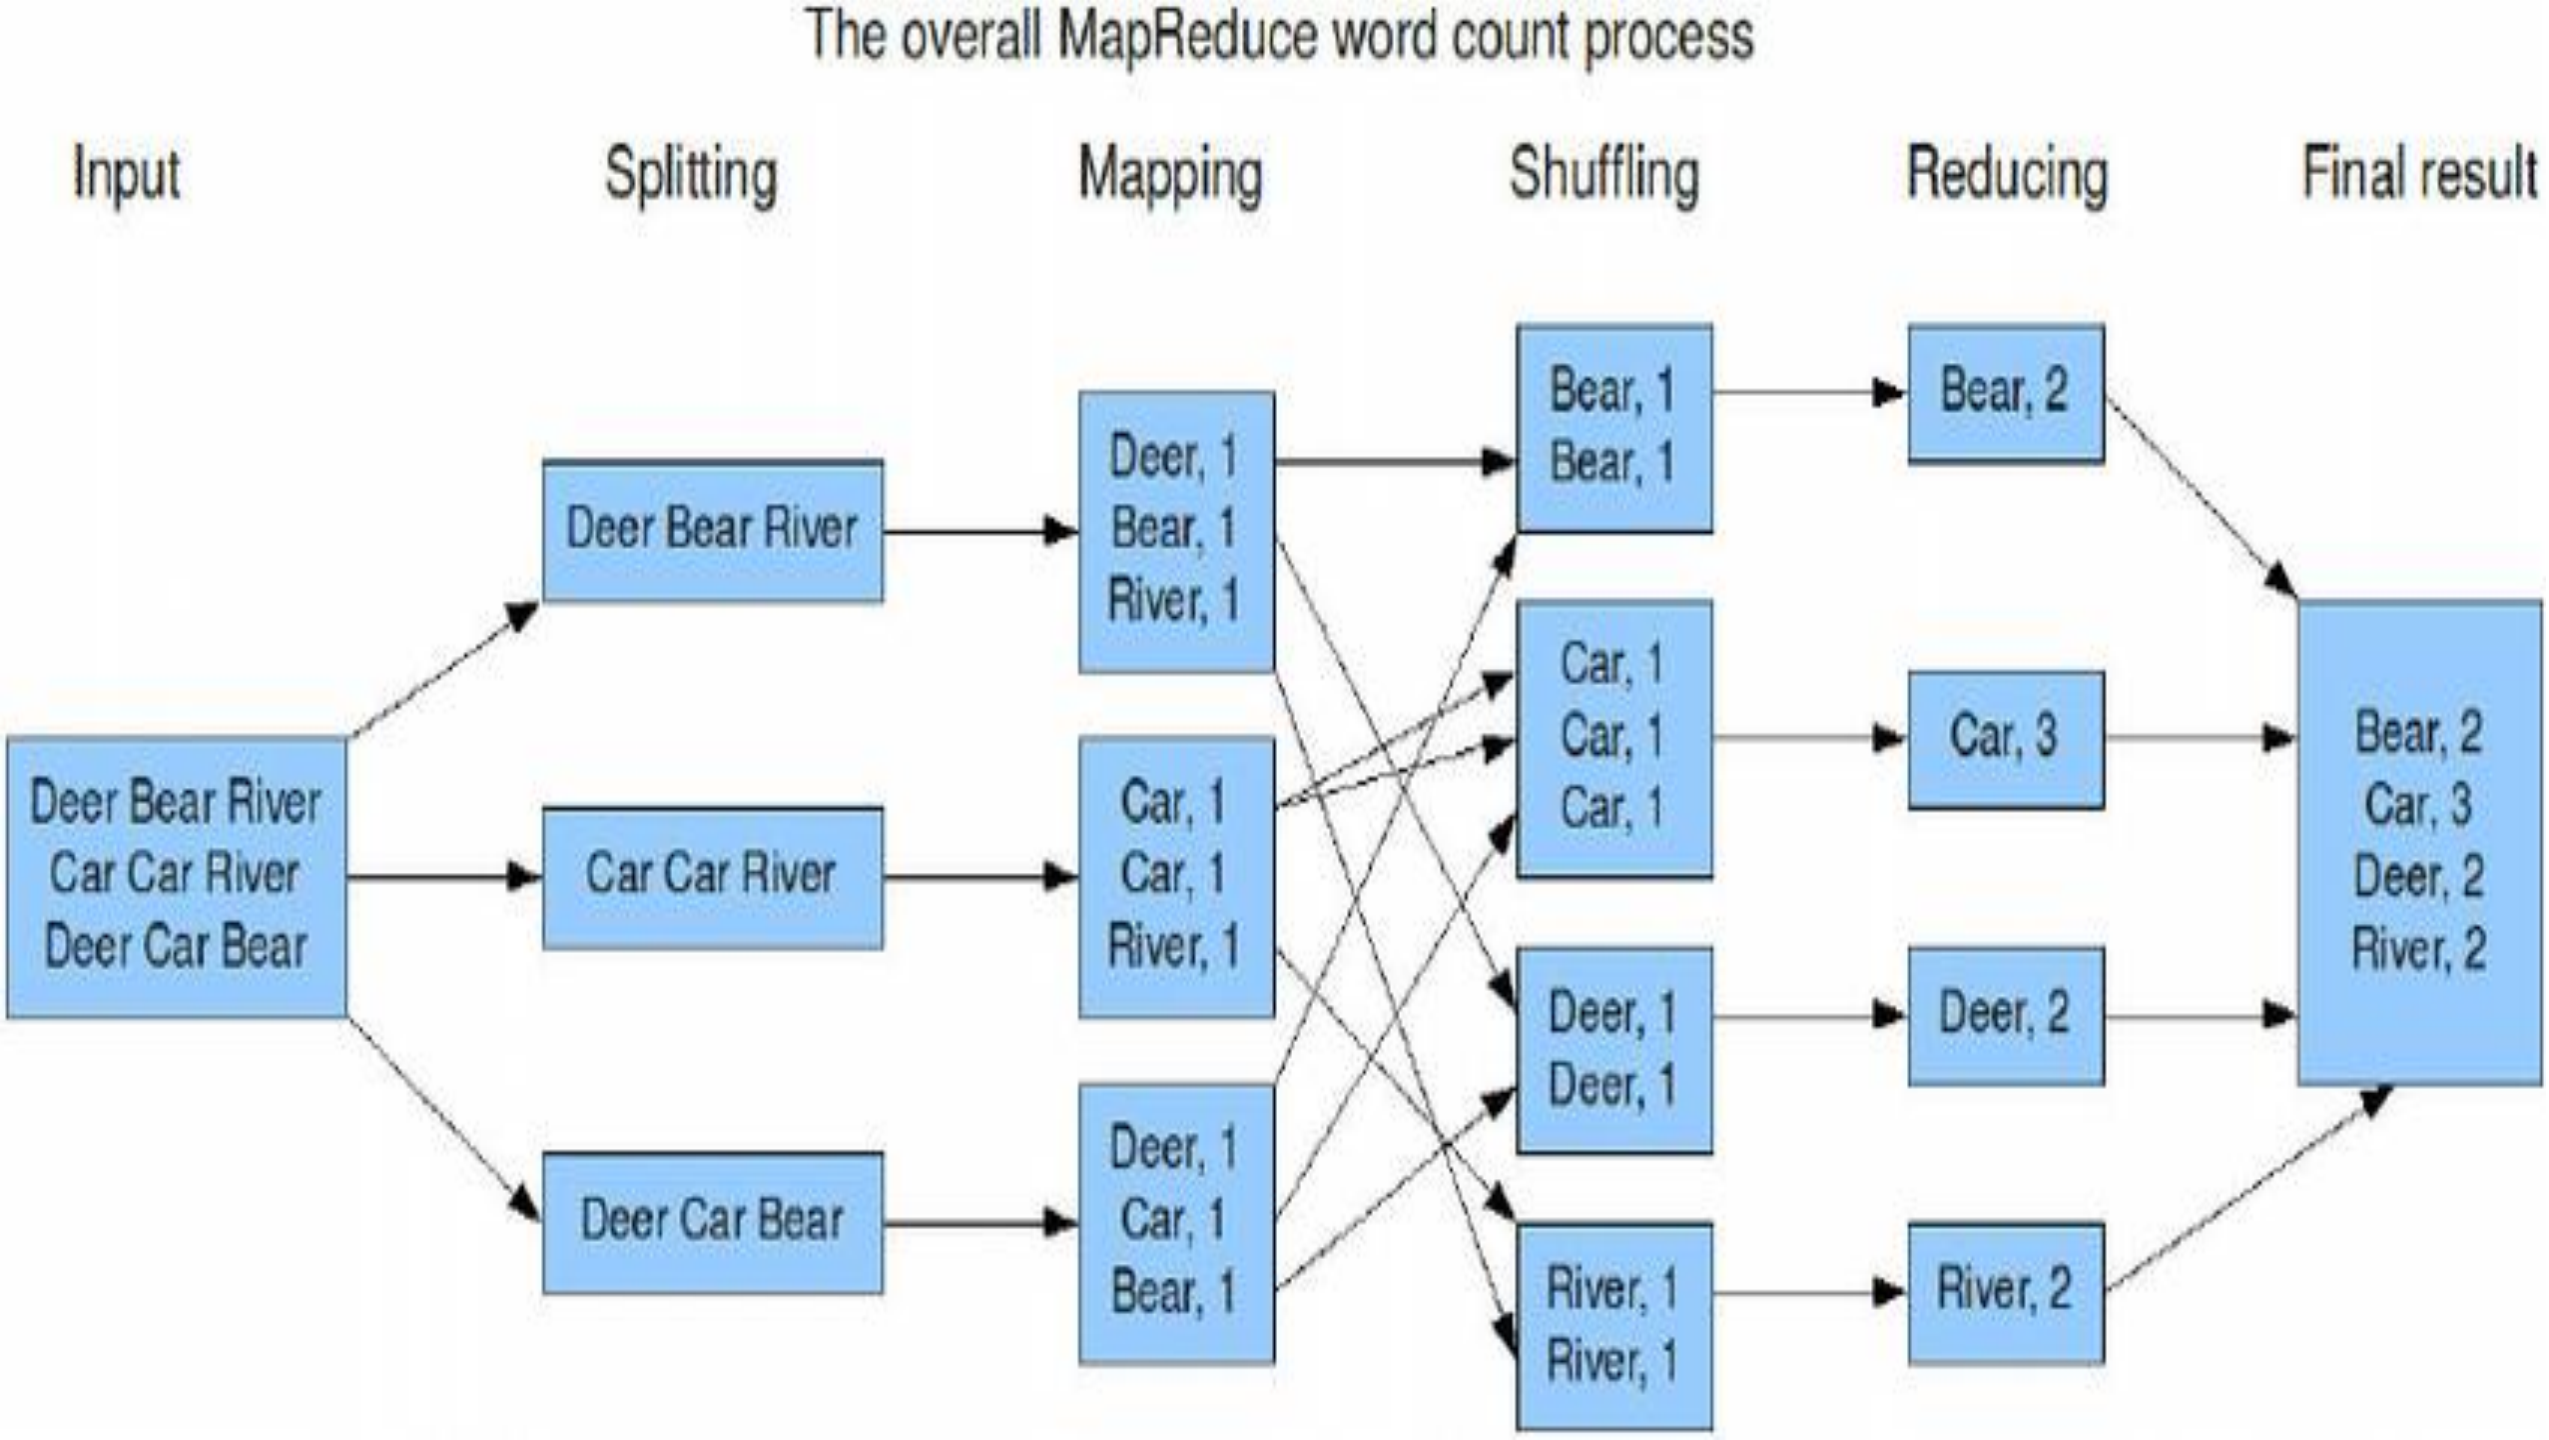
The overall MapReduce word count process
Input
Splitting
Mapping
Shuffling
Reducing
Final result
Deer Bear River
Car Car River
Deer Car Bear
Deer Bear River
Car Car River
Deer Car Bear
Deer, 1
Bear, 1
River, 1
Car, 1
Car, 1
River, 1
Deer, 1
Car, 1
Bear, 1
Bear, 1
Bear, 1
Car, 1
Car, 1
Car, 1
Deer, 1
Deer, 1
River, 1
River, 1
Bear, 2
Car, 3
Deer, 2
River, 2
Bear, 2
Car, 3
Deer, 2
River, 2

MapReduce -> Logical View

- Input and Output data Only from HDFS -> why? -> Java based Distributed storage with Fault Tolerant, Horizontally Scalable

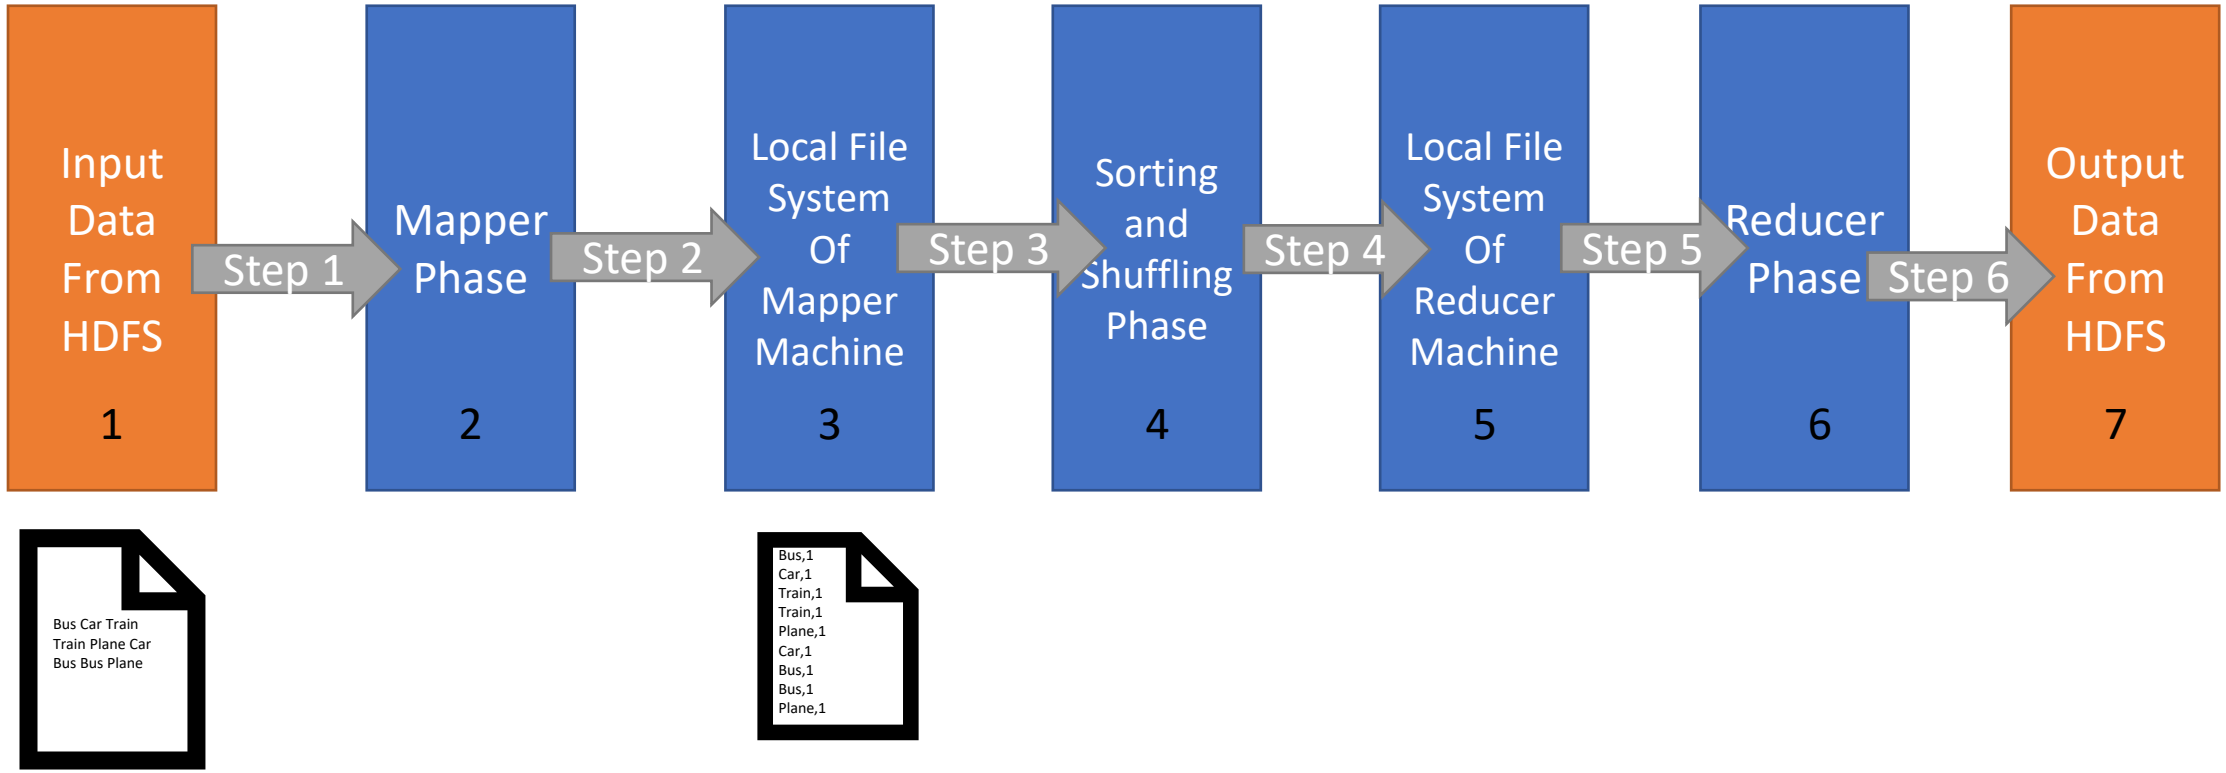

MapReduce -> Physical View

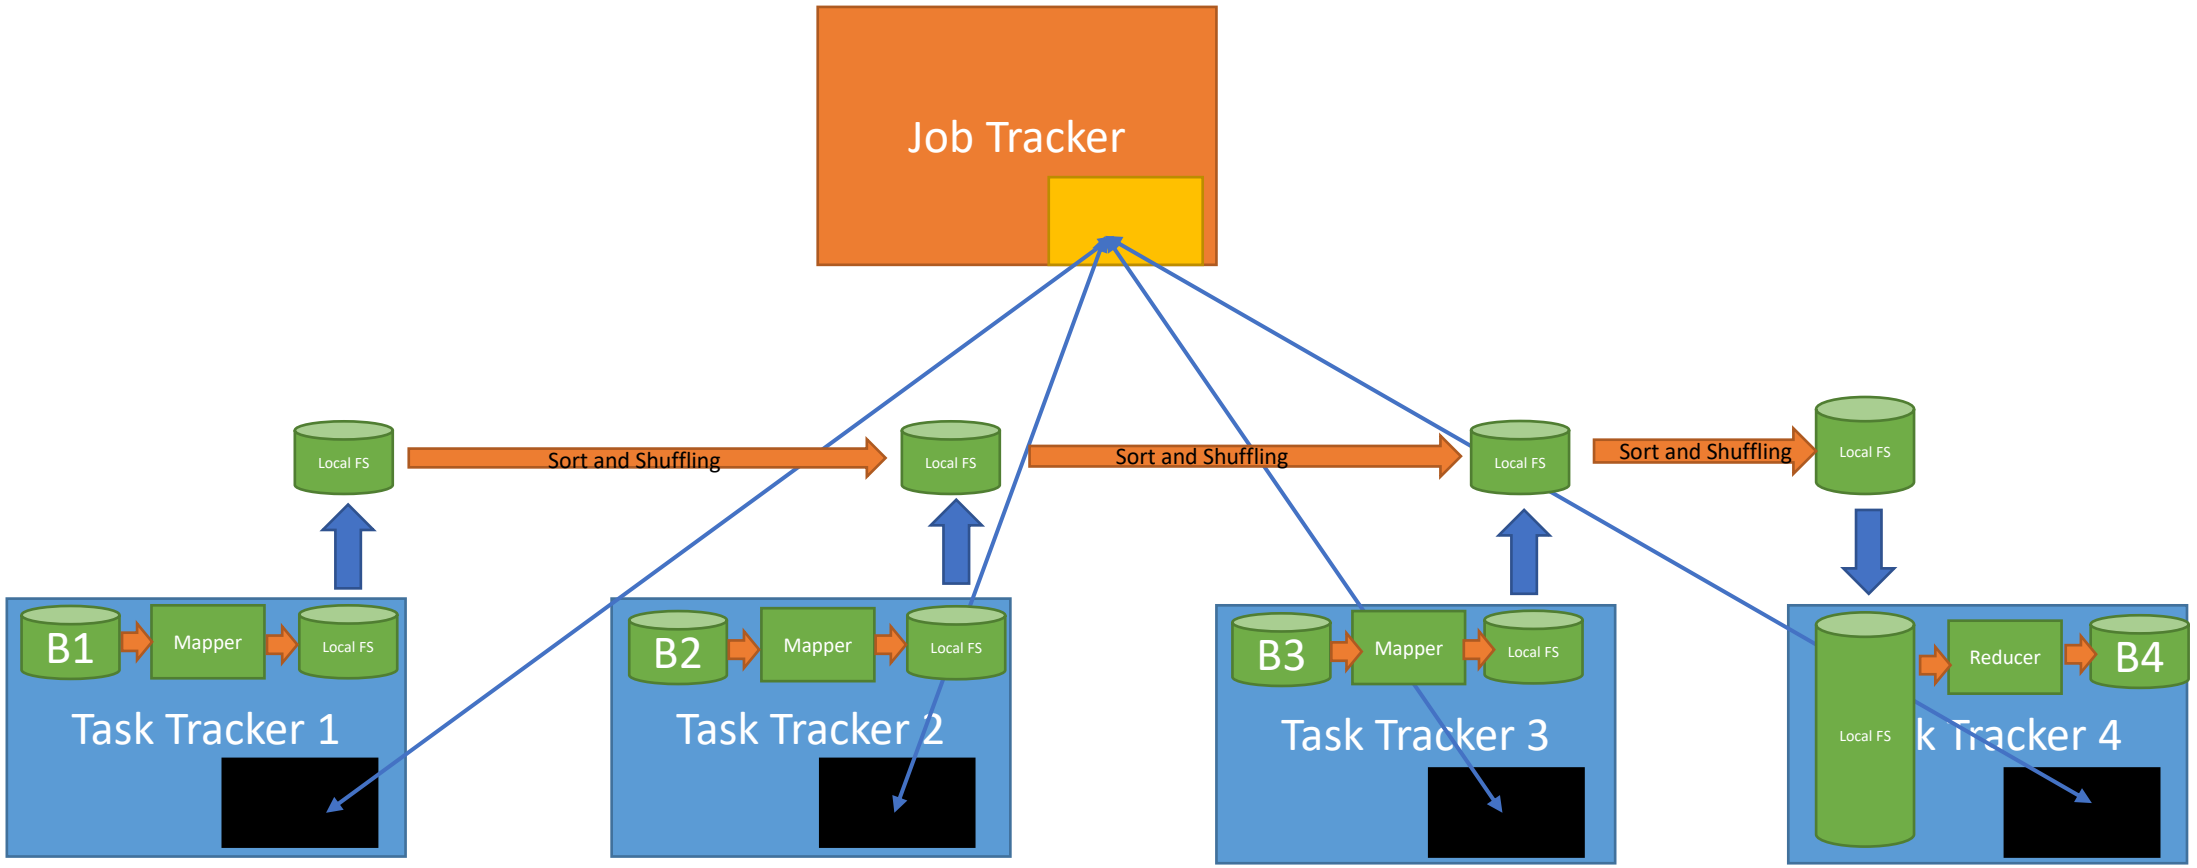

MapReduce -> Physical View

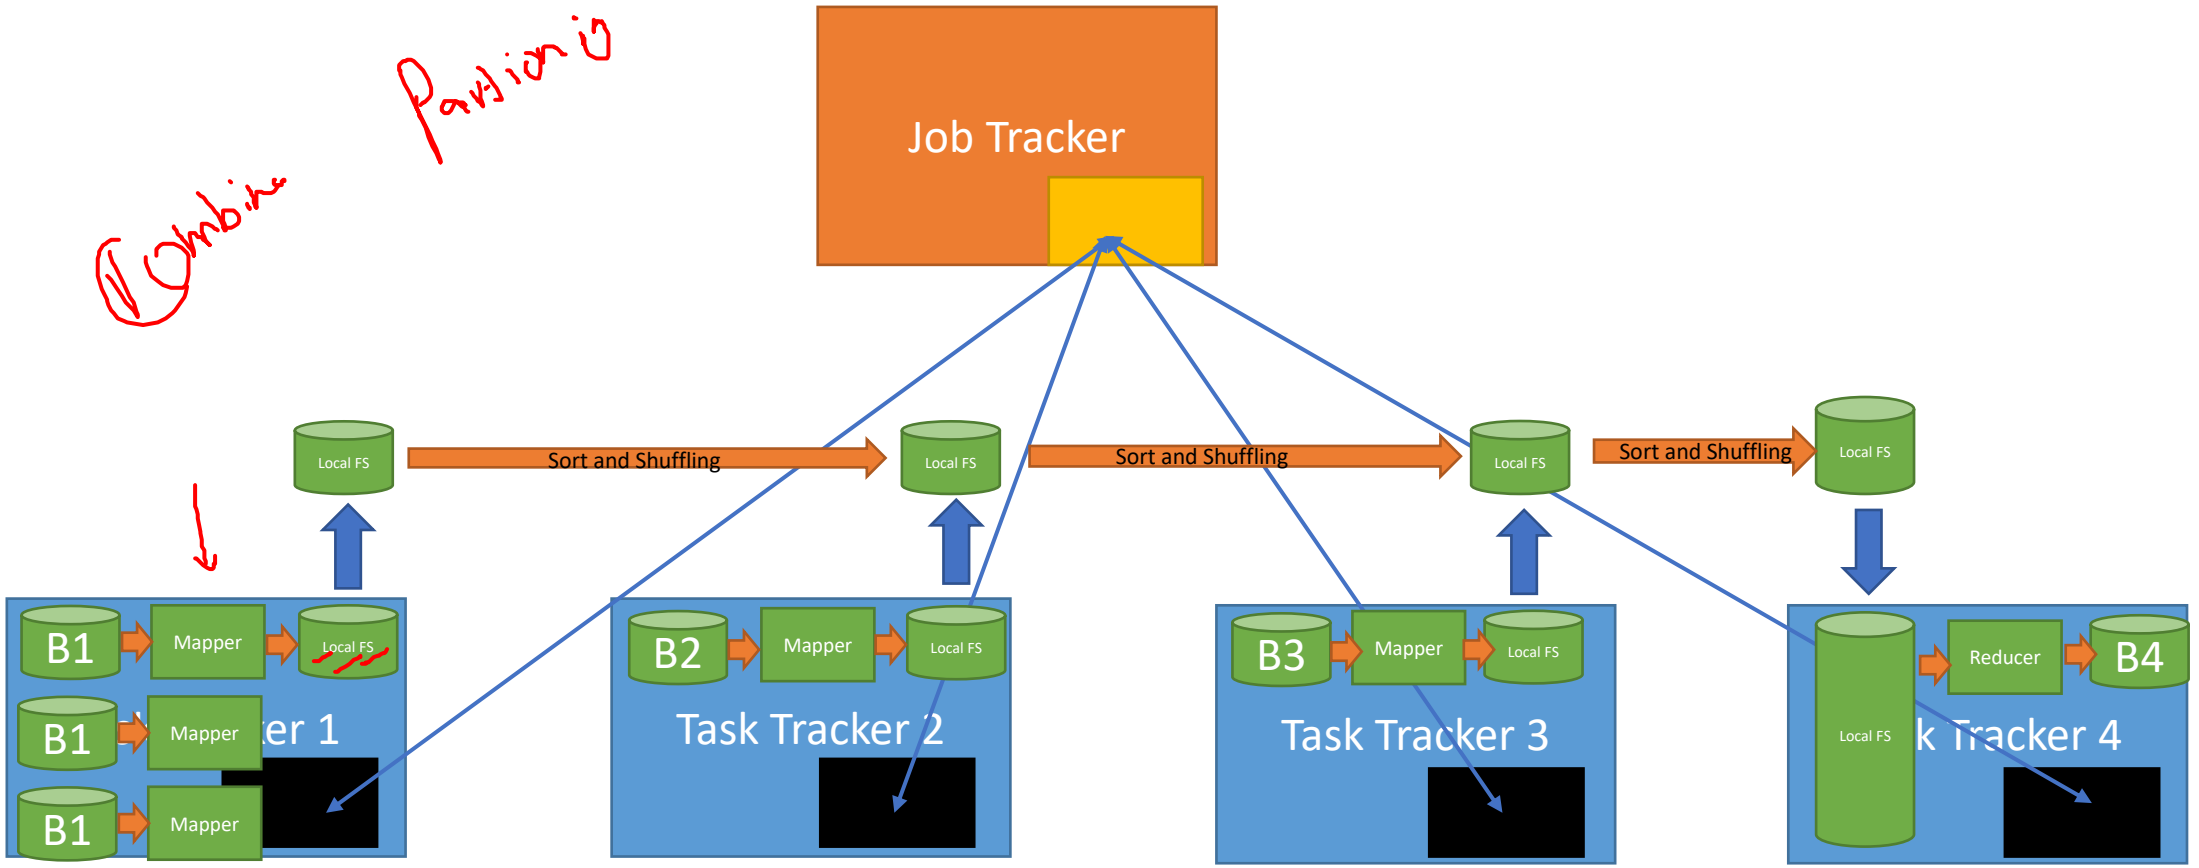

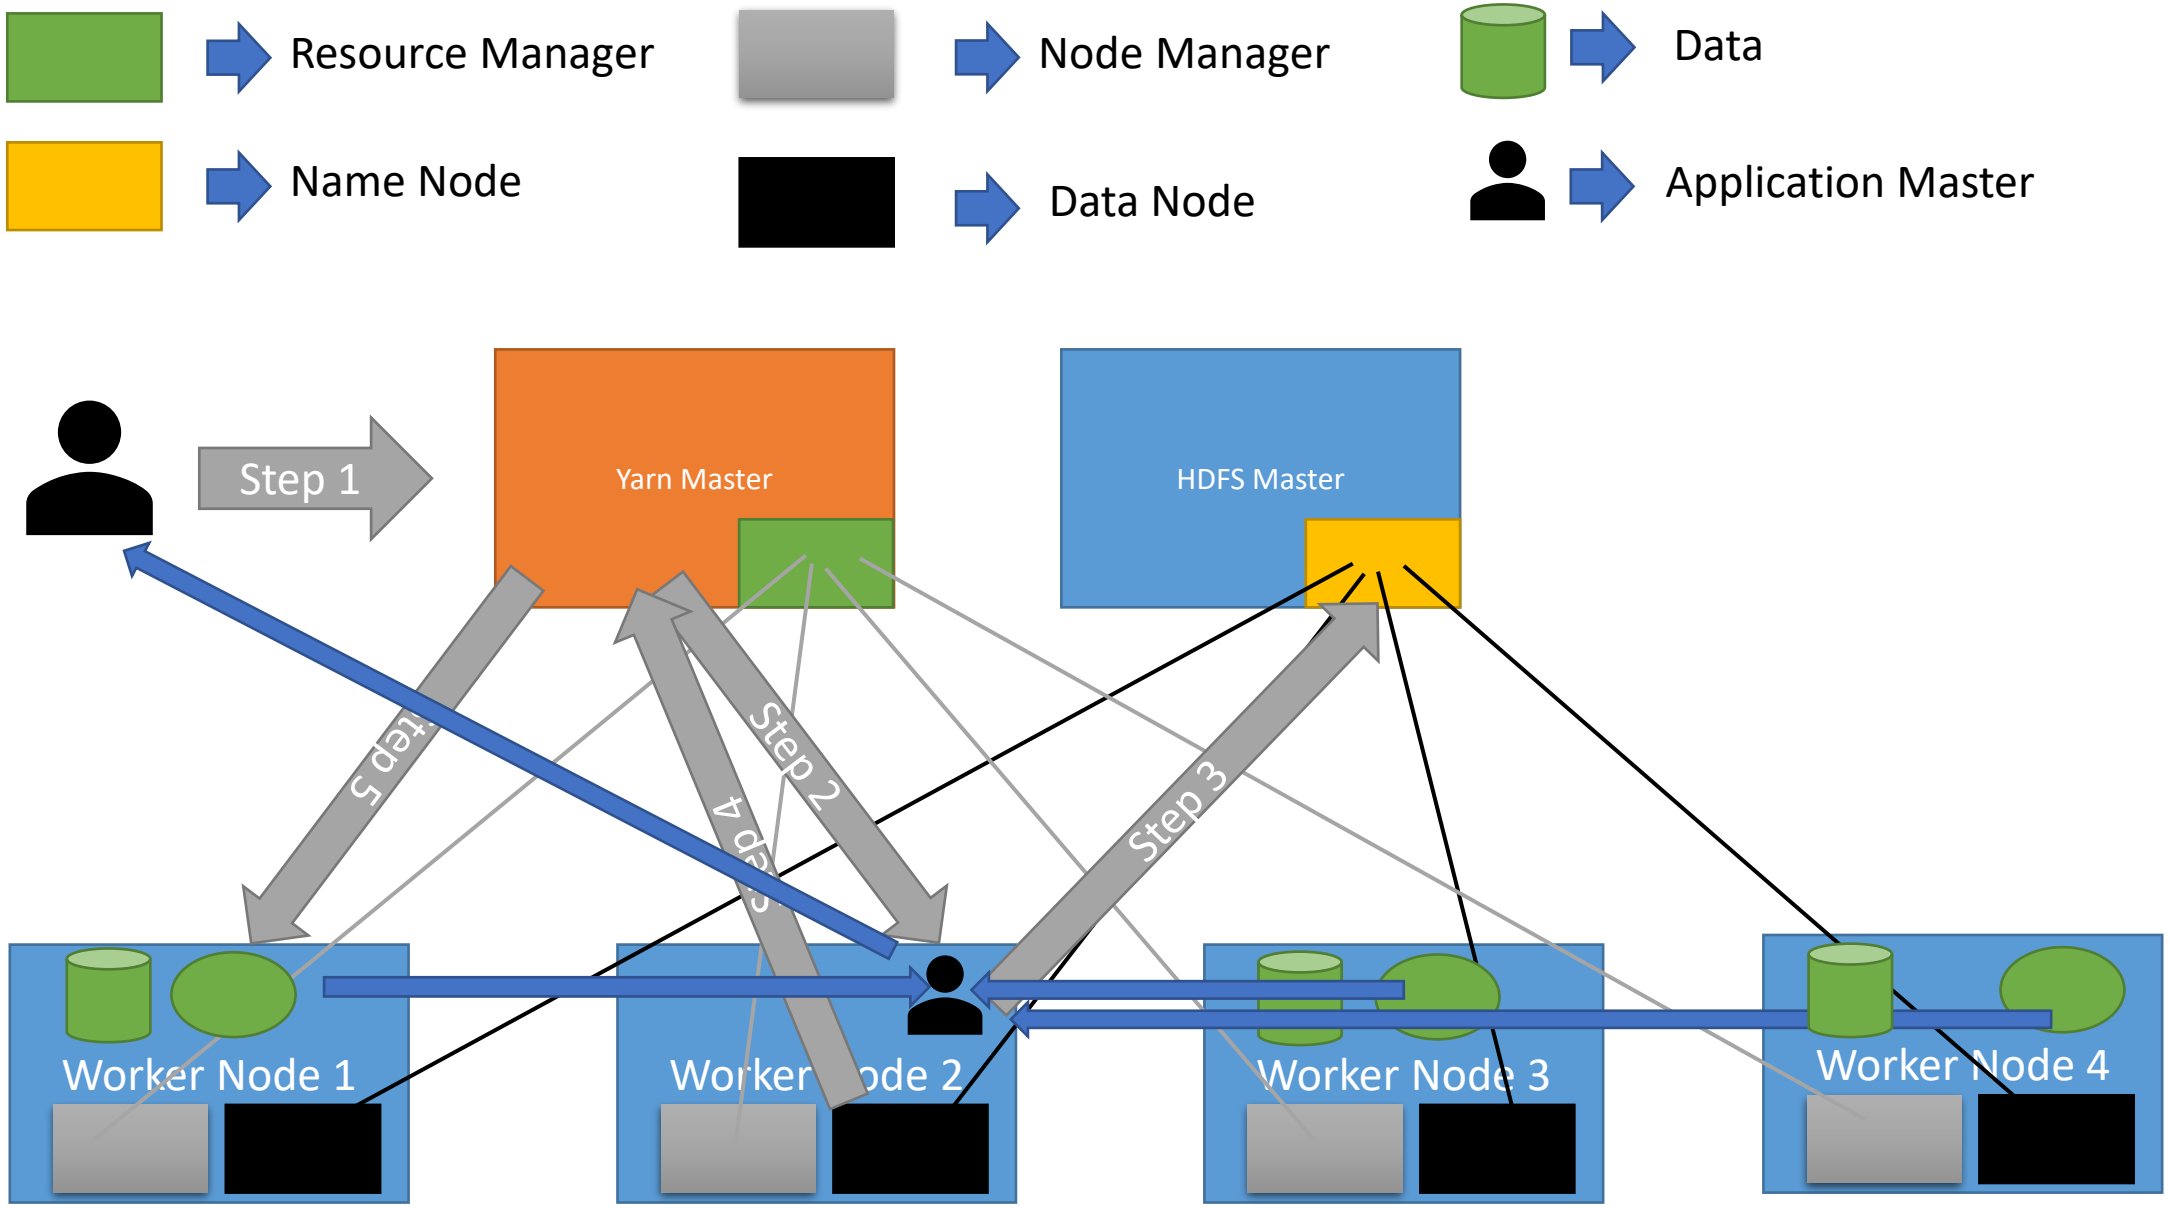
Resource Manager
Node Manager
Data
Name Node
Data Node
Application Master
Step 1
Yarn Master
HDFS Master
Step 2
Step 3
Worker Node 1
Worker Node 2
Worker Node 3
Worker Node 4